

Faculty of Engineering - Ain Shams University (ASUENG)
 Department of Computer and Systems Engineering(CSE)

Course code: CSE352s
 Course Title: Parallel and Distributed Systems
 project : Distributed Online Marketplace System

©2025 David John, Gamal Abdel Shafy Ebrahim. All rights reserved.

\newpage

Team Members

#	Full Name	Role / Contribution Area
1	Youssef Medhat Mahmoud	Architecture Lead, API Gateway & Socket Servers
2	Omar Ahmed Mohamed Ali	Transaction Service & 2PC Protocol
3	Abdallah Hossam Ragab	Database Design & Partitioning Strategy
4	Abdulrahman Ibrahim Mohamed	Authentication Service & 2FA/OTP Module
5	Ahmed Mohamed Mohyeldin	Product Service & Inventory Management
6	Ahmed Ramadan Yassin	Wallet Service & Fund Transfer Logic
7	Ibrahim Mahmoud Ibrahim	SOAP Web Service & XML Protocol Implementation
8	Mostafa Ahmed Abd El Sattar	REST Web Service & JSON Protocol Implementation
9	Tarek Hossam Eldin	Database Replication (Master-Replica Setup)
10	Abdelrahman Ashour Hassan	Frontend UI (Thymeleaf Templates & CSS)
11	Mohamed Hamdy Youssif Ahmed	Testing, QA & Integration Verification
12	Ahmed Ayman Ismael	DevOps, Docker Infrastructure & Deployment

\newpage

Table of Contents

TABLE_OF_CONTENTS_PLACEHOLDER

\newpage

Section 1: Administrative Details

Document Conventions

Throughout this document, the following conventions are used:

- Code snippets are presented in monospaced font within fenced code blocks.
 - Mermaid.js diagrams are embedded directly and can be rendered using the Mermaid Live Editor, Notion, GitHub, or any compatible renderer.
 - Filenames and class names are rendered in monospace.
 - Architectural components are referred to by their service name (e.g., Auth Server, Transaction Server).
 - All monetary values within the system are stored as cents (integer arithmetic) to eliminate floating-point precision errors.
-
-

Section 2: Introduction & Target Beneficiaries

2.1 Introduction to Distributed Online Marketplaces

2.1.1 The Evolution of E-Commerce Architectures

The landscape of electronic commerce has undergone a profound transformation over the past three decades. What began as simple, monolithic web applications serving static product catalogs has evolved into sophisticated distributed systems that must handle millions of concurrent users, process real-time financial transactions, maintain strong consistency guarantees across geographically dispersed data stores, and provide sub-second response times to an increasingly demanding global user base.

The earliest generation of e-commerce platforms -- exemplified by the original architectures of Amazon (circa 1995) and eBay (circa 1995) -- were fundamentally monolithic in nature. A single application binary contained all business logic: user authentication, product catalog management, shopping cart operations, payment processing, and order fulfillment. While this approach offered simplicity in development and deployment, it suffered from critical limitations as user bases scaled:

1. **Vertical Scaling Bottlenecks:** Monolithic architectures require scaling the entire application when any single component experiences load. If the search subsystem receives a spike in traffic, the entire application (including the payment and authentication subsystems) must be replicated, wasting computational resources.
2. **Deployment Coupling:** A minor bug fix in the product catalog module necessitates redeploying the entire application, including unrelated components like the payment processor. This creates risk, increases

downtime, and slows iteration cycles.

3. **Technology Lock-In:** All components within a monolith must share the same technology stack -- the same programming language, the same database, the same framework version. This prevents teams from choosing the most appropriate tool for each specific problem domain.
4. **Fault Isolation Failure:** A memory leak in one module (e.g., the image-processing subsystem) can crash the entire application, taking down all services including the critical payment processing pipeline.

The second generation introduced Service-Oriented Architecture (SOA), which decomposed the monolith into larger, loosely coupled services communicating over standardized protocols such as SOAP (Simple Object Access Protocol) and WSDL (Web Services Description Language). While SOA addressed some of the monolith's limitations, it often resulted in an "ESB (Enterprise Service Bus) bottleneck," where a centralized middleware layer became a single point of failure and a performance constraint.

The current generation -- Microservices Architecture -- represents the natural evolution of SOA principles. Microservices decompose the application into fine-grained, independently deployable services, each owning its own data store, each communicating over lightweight protocols (typically HTTP/REST or gRPC), and each capable of being developed, deployed, and scaled independently. This architecture was popularized by companies like Netflix, Spotify, and Uber, and has become the de facto standard for modern distributed systems.

2.1.2 The Distributed Marketplace Paradigm

An Online Marketplace is a specific class of e-commerce platform that acts as an intermediary between multiple independent sellers and buyers. Unlike a traditional e-commerce store (where a single retailer sells its own inventory), a marketplace enables:

- **Multi-Tenancy:** Multiple sellers list and manage their own products on a shared platform.
- **Peer-to-Peer Transactions:** Financial transfers occur between individual buyers and sellers, mediated by the platform.
- **Decentralized Inventory:** Each seller maintains their own inventory levels; the platform aggregates and presents this information to buyers.
- **Trust Mediation:** The platform provides authentication, transaction guarantees, and dispute resolution mechanisms that enable trust between parties who may never have interacted before.

Designing a marketplace as a distributed system introduces a rich set of computer science challenges that make it an ideal vehicle for studying parallel and distributed systems concepts:

Challenge Domain	Description	Relevant CS Concepts
Concurrent Access	Multiple buyers may attempt to purchase the last unit of ...	Mutual Exclusion, Locking, ACID Transactions
Distributed Data Consistency	User balances, inventory counts, and transaction records ...	CAP Theorem, 2PC, Distributed Transactions
Inter-Process Communication	Autonomous microservices must exchange structured message...	Socket Programming, Protocol Design, TCP/IP
Fault Tolerance	Individual nodes may crash during a multi-step transactio...	Rollback, Compensating Transactions, 2PC
Horizontal Scalability	As user volume grows, individual services must scale inde...	Stateless Services, Database Sharding
Data Partitioning	Large tables must be fragmented across storage nodes to d...	HASH Partitioning, RANGE Partitioning

2.1.3 Project Scope and Objectives

This project implements a Distributed Online Marketplace System that demonstrates the practical application of core parallel and distributed systems concepts within a realistic, production-grade software context. The system is designed

to fulfill the following academic and engineering objectives:

Primary Objectives:

1. Design and implement a distributed microservices architecture using Java 21 and Spring Boot 3.3.5.
2. Build low-level inter-node communication channels using raw Java TCP sockets, manually parsing HTTP request/response frames and constructing JSON/XML payloads without relying on high-level frameworks for the communication layer.
3. Implement a distributed database model using MariaDB 11.4 with native table partitioning (HASH and RANGE strategies) and master-replica replication.
4. Demonstrate atomic distributed transactions using the Two-Phase Commit (2PC) protocol for operations that span multiple logical data stores (wallet debit, inventory decrement, transaction ledger insertion).
5. Provide a complete, functional marketplace UI with Thymeleaf server-side rendering, supporting all CRUD operations, search, purchase, and reporting workflows.

Secondary Objectives:

6. Implement Two-Factor Authentication (2FA) via email-delivered OTP codes for account creation and purchase verification, demonstrating asynchronous messaging patterns.
7. Support both REST and SOAP web service interfaces, implemented at the raw socket level, to demonstrate protocol-level understanding.
8. Enable CSV-based bulk product import for operational efficiency.
9. Provide an External Store API interface for third-party integrations.
10. Implement comprehensive transaction reporting with date-range filtering and user-scoped summary statistics.
11. Integrate an AI Smart Search feature powered by Google Gemini, implemented as a standalone Java socket-based microservice (port 9092), enabling natural language product discovery with relevance scoring and AI-generated search summaries.

2.1.4 Key Differentiators of This Implementation

This project distinguishes itself from typical student marketplace implementations in several critical ways:

1. **Raw Socket HTTP Parsing:** Rather than relying on embedded HTTP servers (Tomcat, Jetty, Netty) for the web service endpoints, the system implements custom `ServerSocket`-based listeners on ports 9090 (REST) and 9091 (SOAP) that manually parse HTTP request lines, headers, and bodies at the byte stream level. This demonstrates a deep understanding of the HTTP protocol that is typically abstracted away by frameworks.
2. **Genuine Database Partitioning:** The MariaDB schema uses native `PARTITION BY HASH` and `PARTITION BY RANGE` directives at the DDL level, distributing data across 4 physical partitions per table. This is not a simulation -- the database engine genuinely distributes rows across separate storage segments, enabling the system to demonstrate partition-aware query routing.
3. **Master-Replica Replication:** The Docker Compose infrastructure provisions two separate MariaDB instances (master on port 3306, read-only replica on port 3307) with ROW-based binary log replication. The Spring Boot application is configured with dual `DataSource` instances (via HikariCP connection pools), enabling read/write splitting in the application layer.
4. **Production-Grade Financial Logic:** All monetary values are stored as `BIGINT` cents (not floating-point), fund transfers include explicit debit-then-credit sequencing with compensating rollback logic, and every transaction generates a cryptographically unique reference code.

2.2 Target Beneficiaries

The Distributed Online Marketplace System is designed to serve four distinct stakeholder categories, each with specific needs, interaction patterns, and value propositions. This section provides an exhaustive analysis of each beneficiary group.

2.2.1 Buyers (End Consumers)

Definition: Buyers are registered users of the marketplace platform who browse, search for, and purchase products listed by sellers. They represent the demand side of the marketplace's two-sided network.

Functional Capabilities Provided to Buyers:

Capability	Description	Implementation Reference
Account Regi	Create a new account with username, email, password, and ...	AuthController.register(), UserService.register()
Secure Authen	Login with BCrypt-hashed password verification. Session-b...	AuthController.login(), UserService.authenticate()
Product Sear	Full-text search across item name, brand, and category fi...	MarketplaceController.searchPage(), ItemService.searchIte...
Product Brow	Browse all active marketplace listings with real-time inv...	MarketplaceController.searchPage(), InventoryService.getA...
Secure Purch	Execute purchases with 2FA OTP verification. System perfo...	MarketplaceController.executePurchase(), TransactionServi...
Wallet Manag	View current balance, deposit funds, and review complete ...	WalletController, WalletService.deposit()
Transaction H	View complete purchase history with date filtering, trans...	ReportController, TransactionService.getBuyerTransactions()
Dashboard O	Centralized dashboard displaying wallet balance, recent p...	DashboardController

Non-Functional Benefits to Buyers:

- **Data Integrity:** The 2PC-style transactional purchase flow ensures that buyers are never charged without receiving the item, and that partial failures (e.g., insufficient seller wallet) result in complete rollback.
- **Real-Time Inventory:** Inventory counts are live-updated with reservation semantics, preventing overselling.
- **Audit Trail:** Every financial operation (deposit, purchase) generates a unique reference code (DEP-XXXXXXXX or TXN-XXXXXXXX) for dispute resolution and record-keeping.

2.2.2 Sellers (Product Listers)

Definition: Sellers are registered users who list products for sale on the marketplace. In this system, any registered user can act as both a buyer and a seller -- there is no distinct seller registration process, reflecting the peer-to-peer nature of the marketplace.

Functional Capabilities Provided to Sellers:

Capability	Description	Implementation Reference
Item Listing (Create)	Add new products with name, description, brand, category,...	ItemController.addItem(), ItemService.createItem()
Item Editing (Update)	Modify existing item details (name, description, brand, c...	ItemController.editItem(), ItemService.updateItem()
Item Removal (Soft Delete)	Remove items from the marketplace. Uses soft-delete (stat...	ItemController.removeItem(), ItemService.removeItem()
Inventory Management	View and update stock quantities for listed items. The sy...	ItemController, InventoryService.updateQuantity()
CSV Bulk Import	Upload a CSV file to create multiple product listings sim...	ItemController.importCsv(), CsvImportService
Sales Dashboard	View completed sales transactions, total revenue, and buy...	DashboardController, TransactionService.getSellerTransact...
Sales Reports	Generate filtered reports of sales activity by date range...	ReportController, ReportService.getTransactionsByDateRan...

Non-Functional Benefits to Sellers:

- **Inventory Reservation:** When a buyer initiates a purchase, the requested quantity is immediately reserved (incremented in the reserved column), preventing other buyers from purchasing stock that is mid-transaction.
- **Automatic Status Management:** When inventory reaches zero, the item status automatically transitions from ACTIVE to SOLD, removing it from search results.
- **Instant Fund Receipt:** Upon completed purchase, the seller's wallet balance is credited immediately (within the same database transaction), with no settlement delay.

2.2.3 External Store Owners (Third-Party Integrators)

Definition: External Store Owners are third-party businesses that integrate with the marketplace through its API interfaces to list products, check inventory, or retrieve account information programmatically.

Integration Points:

Integration Method	Port	Protocol	Format	Description
REST Socket API	9090	HTTP/JSON	JSON	Raw socket-based REST endpoints for item search, account ...
SOAP Socket API	9091	HTTP/XML	SOAP/XML	Raw socket-based SOAP endpoints for transaction reports, ...
Spring REST API	8080	HTTP/JSON	JSON	Framework-managed REST endpoints (via @RestController ann...
Spring SOAP API	8080	HTTP/XML	SOAP/XML	Framework-managed SOAP endpoints (via Spring Web Services)

Authentication for External Stores: The system provides an external_stores table with API key-based authentication. Each external store receives a unique api_key (e.g., ts-api-key-abc123) that must be presented in API requests. The store_listings junction table tracks which marketplace items are cross-listed on external storefronts.

Benefits to External Store Owners:

- **Dual Protocol Support:** Integrators can choose between REST/JSON (modern, lightweight) and SOAP/XML (enterprise, contract-first) depending on their existing infrastructure.
- **Raw Socket Endpoints:** The socket-based APIs on ports 9090 and 9091 demonstrate that the marketplace's inter-node communication is protocol-agnostic and does not depend on framework magic -- integrators see the exact same HTTP frames that internal microservices exchange.
- **Inventory Synchronization:** External stores can poll inventory levels via GET /api/v1/inventory/{itemId} and update stock via PUT /api/v1/inventory/{itemId} to maintain consistency across platforms.

2.2.4 System Administrators

Definition: System Administrators are technical operators responsible for deploying, configuring, monitoring, and

maintaining the marketplace infrastructure.

Administrative Capabilities:

Capability	Description	Tool / Method
Infrastructure Deplo	Provision the entire database infrastructure using Docker...	docker compose up -d
Database Schema	Initialize or migrate the database schema via versioned S...	schema.sql, data.sql
Replication Monitor	Verify master-replica replication health using dedicated ...	replication/verify-replication.sql
Connection Pool M	Monitor HikariCP connection pool metrics for both master ...	application.yml HikariCP configuration
Log Analysis	Review application logs (SLF4J/Logback) for error diagnos...	app.log, logback.xml configuration
Port Management	Manage three separate server processes: Spring Boot (8080...	Process management, netstat diagnostics
Environment Config	Configure SMTP credentials, database URLs, and other secr...	.env, application.yml with \${VAR} syntax
Database Direct Ac	Connect directly to MariaDB instances for ad-hoc queries,...	docker exec -it marketplace-db-master mariadb ...

Non-Functional Benefits to Administrators:

- **Container Isolation:** The Docker-based deployment ensures that MariaDB instances are fully isolated from the host operating system, with persistent volumes for data durability across container restarts.
- **Health Checks:** Both MariaDB containers include Docker health checks (mysqladmin ping) that prevent dependent services from starting before the database is ready.
- **Read-Write Splitting:** The dual-datasource configuration allows administrators to direct read-heavy workloads (search, reports) to the replica, reducing load on the master node.

2.3 Academic Context and Learning Outcomes

This project directly addresses the following learning objectives from the CSE352s curriculum:

Learning Objective	How This Project Addresses It
Understand inter-process comm	Raw TCP socket servers (RestSocketServer, SoapSocketServe...
Design distributed database sc	MariaDB schema uses PARTITION BY HASH and PARTITION BY RA...
Implement distributed transacti	The purchase flow in TransactionService.purchaseItem() im...
Apply concurrency control mec	Spring's @Transactional annotation, combined with InnoDB'...
Understand replication and faul	Master-replica MariaDB replication with ROW-based binary ...
Design and implement web ser	Both REST (JSON over HTTP) and SOAP (XML over HTTP) are i...

Section 3: System Architecture & Adopted

Technologies

3.1 Architectural Overview: Microservices-Based Distributed Model

3.1.1 Architectural Style: Microservices

The Distributed Online Marketplace adopts a Microservices Architecture -- a structural approach in which a single application is developed as a suite of small, independently deployable services, each running in its own process and communicating via lightweight mechanisms. This architectural style, formalized by Martin Fowler and James Lewis in 2014, is characterized by the following principles, all of which are reflected in this implementation:

Principle 1: Single Responsibility (Service Autonomy)

Each microservice in the marketplace owns a distinct business capability:

Service	Business Capability	Entry Point / Class	Port
API Gateway	HTTP request routing, session management, UI rendering	MarketplaceApplication.java (Spring Boot)	8080
Auth Server	User registration, login, password hashing, 2FA/OTP	AuthController, UserService, TwoFactorService	8080*
Product Server	Item CRUD, search, inventory management, CSV import	ItemController, ItemService, InventoryService	8080*
Transaction Server	Purchase orchestration, fund transfer, reporting	TransactionService, WalletService, ReportService	8080*
REST Socket Node	Raw HTTP/JSON web service interface	RestSocketServer.java	9090
SOAP Socket Node	Raw HTTP/XML SOAP web service interface	SoapSocketServer.java	9091

Note (): In the current deployment model, the Auth Server, Product Server, and Transaction Server are co-hosted within the same Spring Boot JVM process (port 8080) for operational simplicity. However, they are architecturally separated at the code level -- distinct controllers, services, repositories, and entities -- and can be extracted into independent deployment units (separate JARs, separate Docker containers) without code changes. This is a deliberate design decision that balances academic demonstration with deployment practicality.*

Principle 2: Decentralized Data Management

Each logical service owns its own data domain. While all tables currently reside in a single MariaDB instance (distributed via partitioning), they are logically separated:

Data Domain	Tables Owned	Partitioning Strategy
User Domain	users, wallets, otp_codes	HASH by user_id (4 partitions)
Product Domain	items, inventory	HASH by seller_id / item_id (4 partitions)
Transaction Domain	transactions, deposit_ledger	RANGE by month / HASH by user_id
Integration Domain	external_stores, store_listings	No partitioning (low-volume reference data)

Principle 3: Smart Endpoints, Dumb Pipes

Inter-service communication uses simple, direct TCP connections -- either via Spring's internal method calls (for co-hosted services) or via raw TCP sockets (for the REST and SOAP nodes). There is no centralized message bus, ESB, or middleware layer. Each endpoint is responsible for understanding the message format and responding appropriately.

Principle 4: Design for Failure

The `TransactionService.purchaseItem()` method implements a multi-step saga with explicit compensation logic:

```
JAVA
// Step 3: Reserve inventory
boolean reserved = inventoryService.reserveStock(itemId, quantity);
if (!reserved) {
    throw new IllegalStateException("Insufficient stock for item: " + item.getName());
}

// Step 4: Transfer funds from buyer to seller
String ref;
try {
    ref = walletService.transfer(buyerId, item.getSellerId(), totalCents);
} catch (Exception e) {
    // Rollback inventory reservation
    LOG.error("Fund transfer failed, rolling back reservation for item {}", itemId);
    throw e; // Spring @Transactional will rollback the entire DB transaction
}
```

If the fund transfer fails (e.g., seller wallet not found), the entire database transaction is rolled back by Spring's `@Transactional` proxy, which issues a ROLLBACK to MariaDB, undoing both the inventory reservation and any partial wallet updates. This is a form of the Saga Pattern with automatic compensating transactions.

3.1.2 Layered Architecture Within Each Service

Within each microservice, the codebase follows a strict four-layer architecture:

```
+-----+
|           Presentation Layer           |
| (Controllers: AuthController, ItemController, |
| MarketplaceController, WalletController, |
| DashboardController, ReportController) |
+-----+
|           Service Layer                |
| (Business Logic: UserService, ItemService, |
| TransactionService, WalletService, |
| InventoryService, TwoFactorService, |
| EmailService, CsvImportService, ReportService) |
+-----+
|           Repository Layer             |
| (Data Access: UserRepository, ItemRepository, |
| WalletRepository, TransactionRepository, |
| InventoryRepository, OtpRepository, |
| DepositLedgerRepository)              |
+-----+
|           Entity Layer                 |
| (Domain Model: User, Item, Wallet, Transaction, |
| Inventory, OtpCode, DepositLedger, |
| ExternalStore)                        |
+-----+
```

- Presentation Layer: Handles HTTP request mapping, session validation, Thymeleaf model population, and

redirect logic. Controllers never contain business logic -- they delegate entirely to the Service Layer.

- Service Layer: Contains all business rules, validation logic, transaction orchestration, and cross-cutting concerns (logging, error handling). Annotated with `@Service` and `@Transactional` for declarative transaction management.
 - Repository Layer: Spring Data JPA interfaces that extend `JpaRepository`. Custom queries use `@Query` annotations with JPQL. No raw SQL in the application layer -- all database access is mediated by Hibernate/JPA.
 - Entity Layer: JPA entity classes annotated with `@Entity`, `@Table`, `@Id`, etc. These are pure data objects with getter/setter methods, `@Enumerated` fields for type-safe status management, and computed properties (e.g., `getBalanceFormatted()`, `getPriceFormatted()`).
-

3.2 Justification of Technology Choices

3.2.1 Why Java 21 and the JVM

Decision: Use Java 21 (LTS) as the primary programming language and runtime platform.

Justification:

1. Native Threading Model: Java provides first-class support for multi-threading via the `java.lang.Thread` class, synchronized blocks, and the `java.util.concurrent` package. This is essential for a project that requires:
 - Spawning a new thread for each incoming socket connection (`RestSocketServer.handleClient()`)
 - Running background daemon threads for the REST and SOAP socket servers
 - Managing connection pools (HikariCP) that maintain thread-safe resource sharing
2. TCP Socket API: Java's `java.net.ServerSocket` and `java.net.Socket` classes provide a clean, well-documented API for raw TCP communication. The course requirement specifically mandates using "Java Sockets" for inter-node communication -- Java's socket API strikes the right balance between low-level control (byte streams) and developer productivity (buffered readers, charset handling).
3. Platform Independence: The JVM's "write once, run anywhere" guarantee simplifies deployment across team members' heterogeneous development environments (Windows, macOS, Linux).
4. Spring Boot Ecosystem: Spring Boot 3.3.5 provides:
 - Embedded Tomcat for the primary web application (port 8080)
 - Spring Data JPA for type-safe, repository-pattern database access
 - Thymeleaf integration for server-side HTML rendering
 - `@Transactional` support for declarative database transaction management
 - Dependency injection via constructor injection (no field injection, following best practices)
 - Spring Mail for SMTP-based email delivery (OTP codes)
5. Modern Language Features (Java 21): The codebase leverages Java 21 features including:
 - Switch Expressions (line 333 of `RestSocketServer.java`):

JAVA

```
String statusText = switch (statusCode) {
    case 200 -> "OK";
    case 400 -> "Bad Request";
    case 404 -> "Not Found";
    case 500 -> "Internal Server Error";
    default -> "OK";
};
```

- Try-with-resources with effectively final variables (line 92 of RestSocketServer.java):

JAVA

```
try (socket;
    BufferedReader in = new BufferedReader(...);
    OutputStream out = socket.getOutputStream()) {
```

- Optional<T> chaining for null-safe entity lookups throughout the Service Layer.

3.2.2 Why Raw Java Sockets Over High-Level Frameworks

Decision: Implement the REST (port 9090) and SOAP (port 9091) web service endpoints using raw `java.net.ServerSocket` and `java.net.Socket` classes, with manual HTTP request parsing, rather than using embedded HTTP servers (Netty, Jetty), HTTP client libraries (Apache HttpClient, OkHttp), or web service frameworks (JAX-RS, JAX-WS).

This is the most critical architectural decision in the project and requires thorough justification.

Academic Justification:

The CSE352s course curriculum explicitly requires students to demonstrate understanding of inter-process communication at the transport layer. Using a high-level framework like Spring's `@RestController` (which is used for the port 8080 endpoints) abstracts away the following critical concepts:

Concept	What the Framework Hides	What Raw Sockets Expose
TCP Connection Lifecycle	Connection accept, keep-alive, timeout handling	<code>ServerSocket.accept()</code> returns a <code>Socket</code> object representin...
HTTP Request Parsing	The framework's embedded server (Tomcat) parses the reque..	Our code manually reads <code>GET /api/v1/items/search HTTP/1.1...</code>
Header Extraction	<code>@RequestHeader</code> , <code>@RequestParam</code> annotations	Manual string splitting: <code>headerLine.substring(0, colon).t...</code>
Content-Length Body Read	Automatic body deserialization	Manual <code>char[] buf = new char[contentLength]; in.read(buf,...</code>
Query String Parsing	<code>@RequestParam</code> automatic binding	Manual URL splitting, <code>URLDecoder.decode()</code> , key-value pair...
JSON Serialization/Deser	Jackson <code>ObjectMapper</code> with <code>@RequestBody</code> / <code>@ResponseBody</code>	Manual <code>StringBuilder</code> JSON construction: <code>sb.append("\"name...</code>
HTTP Response Framing	<code>ResponseEntity<T></code> automatic serialization	Manual HTTP response header construction with <code>Content-Len...</code>
SOAP Envelope Processi	JAX-WS, Spring-WS automatic marshalling/unmarshalling	Manual XML tag extraction: <code>extractXmlValue(xml, "buyerId")</code>
Routing / URL Pattern Ma	<code>@GetMapping("/api/v1/items/{id}")</code> with <code>@PathVariable</code>	Manual regex matching: <code>path.matches("/api/v1/items/\\\\d+...")</code>

The raw socket implementation in this project demonstrates mastery of every layer in the HTTP communication stack:

```

+-----+
| Application Layer (JSON/XML payload construction) | <- We implement this
+-----+
| HTTP Layer (Request line, headers, body framing) | <- We implement this
+-----+
| TCP Layer (Reliable, ordered byte stream delivery) | <- Java Socket API
+-----+
| IP Layer (Packet routing, addressing) | <- Operating System
+-----+
| Link Layer (Ethernet, Wi-Fi frame transmission) | <- Hardware / NIC
+-----+

```

A typical framework-based implementation only requires the student to work at the topmost layer (writing a `@RestController` method). Our raw socket implementation requires working at layers 1 and 2, with direct interaction with layer 3 via the Java Socket API.

Engineering Justification:

Beyond the academic rationale, the raw socket approach provides several engineering benefits for this specific project:

1. **Protocol Customizability:** The custom HTTP parser can be extended to support non-standard headers (e.g., `X-Marketplace-Session-ID`) without framework configuration.
2. **Minimal Dependencies:** The socket servers have zero external dependencies beyond the JDK -- no Tomcat, no Netty, no Jetty. This reduces the attack surface and simplifies deployment.
3. **Educational Debugging:** Students can set breakpoints at the byte-level (`in.readLine()`) and observe the raw HTTP traffic, which is invisible when using framework abstractions.
4. **Thread-Per-Connection Model:** The explicit `new Thread() -> handleClient(clientSocket)` pattern makes the threading model visible and debuggable, unlike the thread pool abstraction hidden inside Tomcat.

3.2.3 Why MariaDB 11.4

Decision: Use MariaDB 11.4 as the distributed database engine, deployed via Docker containers.

Justification:

Criterion	MariaDB 11.4
Native Partitioning	Full support for PARTITION BY HASH and PARTITION BY RANGE...
Replication	Built-in master-replica replication with ROW-based binary...
ACID Compliance	InnoDB storage engine provides full ACID transactions wit...
JDBC Compatibility	The mariadb-java-client JDBC driver (v3.4.1) provides nat...
Docker Availability	Official Docker Hub images (mariadb:11.4) with configurab...
Performance	Thread pool, query plan caching, and InnoDB buffer pool o...
Open Source	GPL v2 license -- no licensing restrictions for academic ...

3.2.4 Why Thymeleaf for Server-Side Rendering

Decision: Use Thymeleaf as the template engine for the web UI, rather than a client-side JavaScript framework (React, Vue, Angular).

Justification: The project's focus is on backend distributed systems, not frontend engineering. Thymeleaf allows the team

to produce a functional, visually presentable UI with minimal JavaScript, while keeping the presentation logic server-side (inside the Spring Boot controllers). This approach:

- Eliminates the need for a separate frontend build pipeline (webpack, Vite, etc.)
- Keeps all routing logic in the Spring controllers (no client-side routing)
- Allows Thymeleaf templates to directly access Java model attributes (e.g., `th:each="item : ${items}"`)
- Renders complete HTML on the server, which is SEO-friendly and requires no client-side hydration

3.2.5 Complete Technology Stack Summary

Technology	Version	Purpose	License
Java (OpenJDK)	21 (LTS)	Primary programming language and runtime	GPL v2 + CPE
Spring Boot	3.3.5	Application framework, dependency injection, web	Apache 2.0
Spring Data JPA	3.3.5	Object-relational mapping, repository pattern	Apache 2.0
Hibernate ORM	6.5+	JPA implementation, query generation, caching	LGPL 2.1
MariaDB	11.4	Distributed relational database (master + replica)	GPL v2
MariaDB JDBC Connector	3.4.1	JDBC driver for database connectivity	LGPL 2.1
HikariCP	5.1+	High-performance JDBC connection pooling	Apache 2.0
Thymeleaf	3.1+	Server-side HTML template engine	Apache 2.0
Spring Security Crypto	6.3+	BCrypt password hashing	Apache 2.0
Spring Session Core	3.3+	HTTP session management	Apache 2.0
Spring Boot Mail	3.3.5	SMTP email delivery for OTP codes	Apache 2.0
SLF4J + Logback	2.0 / 1.5	Structured application logging	MIT / EPL 1.0
dotenv-java	3.0.0	Environment variable loading from .env files	Apache 2.0
Docker	24+	Container runtime for MariaDB instances	Apache 2.0
Docker Compose	2.20+	Multi-container orchestration	Apache 2.0
Maven	3.9+	Build automation, dependency management	Apache 2.0
Java Sockets (JDK)	21	Raw TCP communication for REST/SOAP endpoints	GPL v2 + CPE

3.3 High-Level Architecture Diagram

The following Mermaid.js diagram illustrates the complete system architecture, showing all major components, their communication protocols, and data flow paths:

!Mermaid Diagram 1

3.3.1 Diagram Explanation

The architecture diagram above reveals several key design decisions:

1. Three Entry Points: The system exposes three distinct network entry points:
 - Port 8080 (API Gateway): Handles all browser-based traffic via embedded Tomcat. Serves HTML

pages (Thymeleaf), processes form submissions, and manages HTTP sessions.

- Port 9090 (REST Socket Server): Raw TCP socket listener that manually parses HTTP requests, routes them to the appropriate service methods, and constructs JSON responses.
 - Port 9091 (SOAP Socket Server): Raw TCP socket listener that manually parses HTTP POST requests containing SOAP XML envelopes, extracts operation parameters, and constructs SOAP XML responses.
2. Shared Service Layer: All three entry points (Tomcat, REST Socket, SOAP Socket) share the same service layer instances via Spring's dependency injection container. When the RestSocketServer calls `itemService.searchItems()`, it invokes the exact same `ItemService` bean that `MarketplaceController` uses. This ensures consistent business logic regardless of the communication protocol used by the client.
 3. Master-Replica Database Topology: The Docker Compose file provisions two MariaDB instances:
 - Master (port 3306): Handles all write operations. Configured with `binlog_format = ROW`, `innodb_flush_log_at_trx_commit = 1`, and `sync_binlog = 1` for maximum durability.
 - Replica (port 3307): Configured as `read_only = 1` and `super_read_only = 1`. Receives all changes from the master via binary log replication with a dedicated replicator user (principle of least privilege).
 4. Logical Service Boundaries: Although co-hosted in a single JVM, the three logical servers (Auth, Product, Transaction) have clear boundaries:
 - Auth Server owns user management and 2FA. It has no dependency on transaction or inventory logic.
 - Product Server owns item CRUD and inventory. It does not directly process financial operations.
 - Transaction Server orchestrates the purchase flow, which is the only operation that crosses all three domains (user/wallet + item/inventory + transaction record).
-

3.4 Deployment Topology

The following diagram shows the physical deployment topology:

!Mermaid Diagram 2

Key observations:

- The JVM process and Docker containers run on the same host machine, communicating via localhost.
 - The MariaDB containers are connected via a Docker bridge network (marketplace-net) for inter-container communication (replication traffic).
 - The JVM connects to MariaDB via published ports (3306, 3307) on localhost.
 - All three socket listeners (8080, 9090, 9091) are bound to the same JVM process, sharing the same heap, thread pool, and Spring application context.
-

3.5 Thread Architecture

The system's threading model is critical for understanding its concurrent behavior:

!Mermaid Diagram 3

Threading Model Explanation:

1. Tomcat Threads: The embedded Tomcat server maintains a thread pool (default 200 threads) that handles all HTTP requests on port 8080. Each incoming browser request is dispatched to an available thread from this pool.
2. Socket Server Threads: The RestSocketServer and SoapSocketServer each start a dedicated listener thread (marked as daemon) during @PostConstruct. When a new TCP connection arrives (serverSocket.accept()), a new handler thread is spawned to process that specific request. This is a classic thread-per-connection model. While not as efficient as NIO/event-loop architectures for high concurrency, it provides clear, debuggable code that aligns with the course's threading pedagogy.
3. HikariCP Pools: Two separate connection pools manage database connections:
 - MasterPool (max 10 connections): For read-write queries against the master node.
 - ReplicaPool (max 10 connections): For read-only queries against the replica node.

When a thread needs a database connection, it borrows one from the pool (blocking if all connections are in use), executes its query, and returns the connection. HikariCP's lock-free design ensures minimal contention even under heavy concurrent load.

4. Daemon Threads: All socket server threads are marked as setDaemon(true), meaning they will not prevent JVM shutdown. When the Spring Boot application is stopped (via Ctrl+C or @PreDestroy), the daemon threads are automatically terminated.

End of Sections 1-3. This document continues in subsequent sections covering Application Level Protocol (Section 4), Distributed Database Design (Section 5), Sequence Diagrams (Section 6), Testing (Section 7), End-User Guide (Section 8), and Appendices (Section 9).

Shall I continue with Sections 4 through 6? These sections will cover the custom HTTP-like parser protocol design, JSON/XML packet structures, the distributed database ERD, the Two-Phase Commit protocol deep dive, and comprehensive sequence diagrams for all major flows (Login, Registration, Purchase with 2PC, Deposit, Add/Edit/Remove Item).

Distributed Online Marketplace System --

Documentation (Sections 4-6)

Section 4: Application Level Protocol (Inter-node Communication)

4.1 Overview: Custom HTTP-Like Parser Built Over TCP Sockets

4.1.1 Protocol Design Philosophy

The Distributed Online Marketplace implements a custom application-level protocol that mimics HTTP/1.1 semantics, built directly on top of raw TCP sockets using the Java `java.net.ServerSocket` and `java.net.Socket` APIs. This protocol is designed to:

1. Demonstrate transport-layer understanding by manually constructing and parsing the entire HTTP frame (request line, headers, body) rather than relying on framework-level abstractions.
2. Support dual payload formats -- JSON for the REST interface (port 9090) and XML/SOAP for the SOAP interface (port 9091) -- over the same underlying TCP transport mechanism.
3. Maintain HTTP/1.1 compatibility so that standard HTTP clients (browsers, curl, Postman, SoapUI) can interact with the socket servers without requiring custom client software.

4.1.2 Protocol Stack Architecture

The protocol stack implemented in this project can be visualized as follows:

```
+-----+
| Layer 5: Application Payload                |
| JSON (REST) or SOAP/XML (SOAP) business-specific content |
| e.g. {"itemId": 1, "name": "MacBook Pro", "priceCents": 249999} |
+-----+
| Layer 4: HTTP Frame                        |
| Request Line: GET /api/v1/items/search?q=macbook HTTP/1.1 |
| Headers: Content-Type, Content-Length, Host, SOAPAction    |
| Body: (for POST/PUT) raw payload bytes |
+-----+
| Layer 3: TCP Stream (java.net.Socket)      |
| Reliable, ordered byte stream over TCP    |
| BufferedReader for reading, OutputStream for writing |
+-----+
| Layer 2: IP / Network Layer (OS kernel)    |
+-----+
| Layer 1: Physical / Data Link (NIC hardware) |
+-----+
```

Unlike a framework-based approach (where Tomcat or Netty handles Layers 3-4 and Jackson handles Layer 5), our implementation manually handles Layers 3, 4, and 5 in application code.

4.2 HTTP Request Parsing Pipeline

The `handleClient(Socket socket)` method in both `RestSocketServer.java` and `SoapSocketServer.java` implements a six-stage parsing pipeline. This section provides an exhaustive, code-level walkthrough of each stage.

4.2.1 Stage 1: TCP Connection Accept and Stream Initialization

When a client connects to port 9090 (REST) or 9091 (SOAP), the `ServerSocket.accept()` method returns a new `Socket` object representing the TCP connection. The server wraps this socket's input and output streams in appropriate Java I/O abstractions:

```
JAVA
private void handleClient(Socket socket) {
    try (socket;
        BufferedReader in = new BufferedReader(
            new InputStreamReader(socket.getInputStream(), StandardCharsets.UTF_8));
        OutputStream out = socket.getOutputStream()) {
        // ... parsing pipeline ...
    }
}
```

Key design decisions:

- `BufferedReader` is used for the input stream because HTTP is a text-based, line-oriented protocol (request line and headers are terminated by `\r\n`).
- `OutputStream` (not `PrintWriter`) is used for the output stream because we need precise control over the byte encoding of the response (including the `Content-Length` header which must match the exact byte count of the body).
- `StandardCharsets.UTF_8` is explicitly specified to ensure consistent encoding across platforms.
- The `try-with-resources` block ensures that the socket and streams are automatically closed when the handler finishes, preventing resource leaks.

4.2.2 Stage 2: Request Line Parsing

The first line of an HTTP request is the request line, which contains the method, URI, and protocol version:

```
GET /api/v1/items/search?q=macbook HTTP/1.1\r\n
```

The parser extracts this line and splits it into its three components:

```
JAVA
// Stage 2: Parse the HTTP request line
String requestLine = in.readLine();
if (requestLine == null || requestLine.isEmpty()) return;

String[] parts = requestLine.split(" ");
if (parts.length < 2) return;
String method = parts[0]; // "GET", "PUT", "POST"
String fullPath = parts[1]; // "/api/v1/items/search?q=macbook"
```

Error handling: If the request line is null (connection closed before sending data) or malformed (fewer than 2

space-separated tokens), the handler returns immediately without sending a response, allowing the socket to close cleanly.

4.2.3 Stage 3: Header Parsing

HTTP headers follow the request line, each on its own line in the format Header-Name: Header-Value. The header section is terminated by an empty line (`\r\n\r\n`):

```
JAVA
// Stage 3: Parse headers
Map<String, String> headers = new HashMap<>();
String headerLine;
int contentLength = 0;
while ((headerLine = in.readLine()) != null && !headerLine.isEmpty()) {
    int colon = headerLine.indexOf(':');
    if (colon > 0) {
        String key = headerLine.substring(0, colon).trim().toLowerCase();
        String value = headerLine.substring(colon + 1).trim();
        headers.put(key, value);
        if (key.equals("content-length")) {
            contentLength = Integer.parseInt(value);
        }
    }
}
```

Implementation details:

- Headers are stored in a case-insensitive manner (keys are lowercased) to comply with HTTP/1.1 RFC 7230, which specifies that header names are case-insensitive.
- The Content-Length header is extracted as a special case because it controls how many bytes of body data to read in the next stage.
- For the SOAP server, the SOAPAction header is also extracted to potentially identify the requested operation (though the actual implementation detects operations from the XML body).

4.2.4 Stage 4: Body Reading

For methods that include a request body (PUT, POST), the body is read based on the Content-Length header value:

```
JAVA
// Stage 4: Read body (for PUT/POST)
String body = "";
if (contentLength > 0) {
    char[] buf = new char[contentLength];
    int totalRead = 0;
    while (totalRead < contentLength) {
        int read = in.read(buf, totalRead, contentLength - totalRead);
        if (read < 0) break;
        totalRead += read;
    }
    body = new String(buf, 0, totalRead);
}
```

Critical note: The SOAP server uses a loop-based read that handles partial TCP reads. Because TCP is a streaming protocol, a single `read()` call may not return all requested bytes -- the data might arrive across multiple TCP segments. The loop continues reading until either `contentLength` bytes have been read or the stream ends.

4.2.5 Stage 5: Query String Parsing and URL Routing

The path and query string are separated at the ? character, and query parameters are decoded:

```

JAVA
// Stage 5: Separate path and query string
String path = fullPath;
Map<String, String> queryParams = new HashMap<>();
int qMark = fullPath.indexOf('?');
if (qMark >= 0) {
    path = fullPath.substring(0, qMark);
    String queryStr = fullPath.substring(qMark + 1);
    for (String pair : queryStr.split("&")) {
        String[] kv = pair.split("=", 2);
        String key = URLDecoder.decode(kv[0], StandardCharsets.UTF_8);
        String value = kv.length > 1
            ? URLDecoder.decode(kv[1], StandardCharsets.UTF_8) : "";
        queryParams.put(key, value);
    }
}

```

The routing logic uses regex matching to dispatch to the appropriate handler:

```

JAVA
if (method.equals("GET") && path.equals("/api/v1/items/search")) {
    responseJson = handleItemSearch(queryParams);
    statusCode = 200;
} else if (method.equals("GET") && path.matches("/api/v1/items/\\d+")) {
    Long itemId = Long.parseLong(path.substring("/api/v1/items/".length()));
    responseJson = handleGetItem(itemId);
    statusCode = responseJson.contains("\"error\"") ? 404 : 200;
}
// ... additional routes ...

```

4.2.6 Stage 6: HTTP Response Construction

The final stage constructs a complete HTTP/1.1 response with status line, headers, and body:

```

JAVA
private void sendHttpResponse(OutputStream out, int statusCode,
                             String contentType, String body) throws IOException {
    String statusText = switch (statusCode) {
        case 200 -> "OK";
        case 400 -> "Bad Request";
        case 404 -> "Not Found";
        case 500 -> "Internal Server Error";
        default -> "OK";
    };

    byte[] bodyBytes = body.getBytes(StandardCharsets.UTF_8);
    String response = "HTTP/1.1 " + statusCode + " " + statusText + "\r\n" +
        "Content-Type: " + contentType + "; charset=utf-8\r\n" +
        "Content-Length: " + bodyBytes.length + "\r\n" +
        "Access-Control-Allow-Origin: *\r\n" +
        "Connection: close\r\n" +
        "\r\n";
    out.write(response.getBytes(StandardCharsets.UTF_8));
    out.write(bodyBytes);
    out.flush();
}

```

Design decisions:

- Content-Length is computed from bodyBytes.length (the byte-level length after UTF-8 encoding), not from body.length() (the character count), which avoids mismatches when the body contains multi-byte UTF-8 characters.
 - Connection: close is set because the server uses a thread-per-connection model without HTTP keep-alive support.
 - Access-Control-Allow-Origin: * enables CORS for browser-based API testing.
-

4.3 JSON Packet Structures (REST Interface -- Port 9090)

This section defines the exact JSON packet structures for all REST endpoints. Each structure shows the request format, response format, and relevant status codes.

4.3.1 Request/Response 1: Item Search

Request:

```
GET /api/v1/items/search?q=macbook&excludeSeller=1 HTTP/1.1
Host: localhost:9090
Accept: application/json
Connection: close
```

Response (200 OK):

```
JSON
{
  "query": "macbook",
  "resultCount": 1,
  "items": [
    {
      "itemId": 1,
      "name": "MacBook Pro 16\"",
      "description": "Apple MacBook Pro with M3 chip, 16GB RAM",
      "brand": "Apple",
      "category": "Electronics",
      "priceCents": 249999,
      "priceFormatted": "$2,499.99",
      "status": "ACTIVE",
      "sellerId": 1,
      "availableQty": 5
    }
  ]
}
```

Field	Type	Description
query	String	The original search query echoed back
resultCount	Integer	Number of matching items
items[].itemId	Long	Unique item identifier
items[].priceCents	Long	Price in cents (integer arithmetic)
items[].priceFormatted	String	Human-readable price string
items[].availableQty	Integer	quantity - reserved from inventory

4.3.2 Request/Response 2: Account Information

Request:

```
GET /api/v1/accounts/1 HTTP/1.1
Host: localhost:9090
Accept: application/json
```

Response (200 OK):

```
JSON
{
  "userId": 1,
  "username": "alice",
  "fullName": "Alice Johnson",
  "email": "alice@example.com",
  "balanceCents": 10000000,
  "balanceFormatted": "$100,000.00",
  "itemsForSale": 3,
  "totalPurchases": 5,
  "totalSales": 2
}
```

Response (404 Not Found):

```
JSON
{
  "error": "User not found"
}
```

4.3.3 Request/Response 3: Inventory Query

Request:

```
GET /api/v1/inventory/1 HTTP/1.1
Host: localhost:9090
```

Response (200 OK):

```
JSON
{
  "itemId": 1,
  "quantity": 5,
  "reserved": 0,
  "available": 5
}
```

4.3.4 Request/Response 4: Inventory Update (PUT)

Request:

```
PUT /api/v1/inventory/1 HTTP/1.1
Host: localhost:9090
Content-Type: application/json
Content-Length: 16

{"quantity": 10}
```

Response (200 OK):

```
JSON
{
  "itemId": 1,
  "quantity": 10,
  "available": 10,
  "message": "Inventory updated successfully"
}
```

Response (400 Bad Request):

```
JSON
{
  "error": "Inventory not found for item: 999"
}
```

4.3.5 Request/Response 5: SOAP Purchase Item

Request:

```

POST /ws HTTP/1.1
Host: localhost:9091
Content-Type: text/xml; charset=utf-8
SOAPAction: "purchaseItem"
Content-Length: 412

<?xml version="1.0" encoding="UTF-8"?>
<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:ns="http://marketplace.com/soap">
  <soap:Body>
    <ns:purchaseItemRequest>
      <ns:buyerId>2</ns:buyerId>
      <ns:itemId>1</ns:itemId>
      <ns:quantity>1</ns:quantity>
      <ns:otpCode>123456</ns:otpCode>
    </ns:purchaseItemRequest>
  </soap:Body>
</soap:Envelope>

```

Response (200 OK -- Success):

```

XML
<?xml version="1.0" encoding="UTF-8"?>
<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:ns="http://marketplace.com/soap">
  <soap:Body>
    <ns:purchaseItemResponse>
      <ns:success>true</ns:success>
      <ns:message>Purchase completed successfully</ns:message>
      <ns:transactionId>42</ns:transactionId>
      <ns:referenceCode>TXN-A1B2C3D4</ns:referenceCode>
    </ns:purchaseItemResponse>
  </soap:Body>
</soap:Envelope>

```

Response (200 OK -- Failure):

```

XML
<?xml version="1.0" encoding="UTF-8"?>
<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:ns="http://marketplace.com/soap">
  <soap:Body>
    <ns:purchaseItemResponse>
      <ns:success>false</ns:success>
      <ns:message>Invalid or expired OTP code</ns:message>
    </ns:purchaseItemResponse>
  </soap:Body>
</soap:Envelope>

```

4.3.6 Custom Status Code Reference

Status Code	Status Text	Usage in System
200	OK	Successful GET, successful PUT, successful SOAP operation
400	Bad Request	Invalid input (e.g., malformed JSON body, business rule v...
404	Not Found	Resource not found (e.g., unknown item ID, user ID, unmat...
500	Internal Server Error	Unhandled exception in handler logic

4.4 Sequence Diagram: Byte-Level Socket Communication

The following Mermaid.js sequence diagram shows the complete byte-level communication handshake and data exchange between an HTTP client and the REST Socket Server for an item search request:

```
sequenceDiagram
    participant Client as HTTP Client<br/>(curl/Browser)
    participant TCP as TCP Layer<br/>(OS Kernel)
    participant Server as RestSocketServer<br/>(Port 9090)
    participant Thread as Handler Thread<br/>(Daemon)
    participant Service as ItemService<br/>(Spring Bean)
    participant DB as MariaDB<br/>(Master:3306)

    Note over Client,TCP: TCP Three-Way Handshake
    Client->>TCP: SYN (port 9090)
    TCP->>Server: SYN-ACK
    Client->>TCP: ACK
    Note over Client,TCP: TCP Connection Established

    Note over Server,Thread: ServerSocket.accept() returns Socket
    Server->>Thread: new Thread(handleClient(socket))

    Note over Client,Thread: HTTP Request (Byte Stream)
    Client->>Thread: "GET /api/v1/items/search?q=macbook HTTP/1.1\r\n"
    Client->>Thread: "Host: localhost:9090\r\n"
    Client->>Thread: "Accept: application/json\r\n"
    Client->>Thread: "\r\n" (empty line = end of headers)

    Note over Thread: Stage 1: in.readLine() -> request line
    Note over Thread: Stage 2: Parse method="GET", path="/api/v1/items/search?q=macbook"
    Note over Thread: Stage 3: Read headers until empty line
    Note over Thread: Stage 4: Content-Length=0, skip body
    Note over Thread: Stage 5: Split path at '?', decode query params

    Thread->>Service: itemService.searchItems("macbook", 0)
    Service->>DB: SELECT * FROM items WHERE status='ACTIVE'<br/>AND (name LIKE '%macbook%'<br/>OR brand LIK...
    DB-->>Service: ResultSet (1 row)
    Service-->>Thread: List<Item> [MacBook Pro]

    Note over Thread: Stage 6: Build JSON with StringBuilder
    Note over Thread: Calculate Content-Length from UTF-8 bytes

    Thread->>Client: "HTTP/1.1 200 OK\r\n"
    Thread->>Client: "Content-Type: application/json; charset=utf-8\r\n"
    Thread->>Client: "Content-Length: 287\r\n"
    Thread->>Client: "Access-Control-Allow-Origin: *\r\n"
    Thread->>Client: "Connection: close\r\n"
    Thread->>Client: "\r\n"
    Thread->>Client: {"query":"macbook","resultCount":1,"items":[...]}}

    Note over Client,TCP: TCP Connection Close
    Thread->>TCP: socket.close()
    TCP->>Client: FIN
    Client->>TCP: FIN-ACK
```

4.4.1 Diagram Explanation

The sequence diagram above reveals the complete lifecycle of a single REST API request:

1. TCP Handshake (Lines 1-3): The standard three-way TCP handshake (SYN, SYN-ACK, ACK) establishes a reliable byte stream connection between the client and port 9090.
2. Thread Spawn (Line 4): The `ServerSocket.accept()` call in the main listener thread returns a new `Socket`

object. A new daemon thread is spawned to handle this specific connection, allowing the listener to immediately return to accepting new connections.

3. HTTP Request Transmission (Lines 5-8): The client sends the HTTP request as a sequence of ASCII text lines over the TCP byte stream. Each line is terminated by `\r\n` (carriage return + line feed). The empty line (`\r\n\r\n`) signals the end of the header section.
4. Parsing Pipeline (Lines 9-13): The handler thread executes the six-stage parsing pipeline described in Section 4.2, extracting the method, path, query parameters, and (if applicable) request body.
5. Service Invocation (Lines 14-16): The handler calls into the Spring-managed `ItemService` bean, which in turn executes a JPQL query via Hibernate, which generates a SQL `SELECT` statement against the MariaDB master node.
6. Response Construction (Lines 17-24): The handler manually constructs the HTTP response, calculating the exact byte length of the JSON body for the Content-Length header, then writes the response headers and body to the output stream.
7. Connection Teardown (Lines 25-27): The handler closes the socket, triggering a TCP FIN exchange that gracefully terminates the connection.

4.5 SOAP Envelope Processing

The SOAP Socket Server (port 9091) implements the same TCP/HTTP parsing pipeline as the REST server, but with an additional layer of XML processing on top:

4.5.1 XML Tag Extraction

Since the project avoids XML parsing libraries (DOM, SAX, StAX), XML values are extracted using a custom string-matching algorithm:

JAVA

```
private String extractXmlValue(String xml, String tagName) {
    // Try namespace prefixes: "ns:", "soap:", "" (none)
    String[] prefixes = { "ns:", "soap:", "" };
    for (String prefix : prefixes) {
        String openTag = "<" + prefix + tagName + ">";
        String closeTag = "</" + prefix + tagName + ">";
        int start = xml.indexOf(openTag);
        if (start >= 0) {
            start += openTag.length();
            int end = xml.indexOf(closeTag, start);
            if (end >= 0) return xml.substring(start, end).trim();
        }
    }
    // Fallback: try any namespace prefix
    int tagStart = xml.indexOf(": " + tagName + ">");
    if (tagStart >= 0) {
        int contentStart = xml.indexOf(">", tagStart) + 1;
        int contentEnd = xml.indexOf("<", contentStart);
        if (contentEnd > contentStart)
            return xml.substring(contentStart, contentEnd).trim();
    }
    return "";
}
```

This approach handles three common XML namespace patterns:

- <ns:buyerId>2</ns:buyerId> (explicit namespace prefix)
- <soap:Body> (SOAP envelope namespace)
- <buyerId>2</buyerId> (default namespace, no prefix)

4.5.2 SOAP Operation Detection

Rather than relying on the SOAPAction HTTP header (which is optional and not always present), the server detects operations by inspecting the XML body content:

```
JAVA
if (xmlBody.contains("getTransactionReportRequest")) {
    responseXml = handleGetTransactionReport(xmlBody);
} else if (xmlBody.contains("purchaseItemRequest")) {
    responseXml = handlePurchaseItem(xmlBody);
} else if (xmlBody.contains("getUserInfoRequest")) {
    responseXml = handleGetUserInfo(xmlBody);
} else {
    responseXml = buildSoapFault("Unknown operation");
}
```

This is a pragmatic design choice: rather than implementing a full WS-Addressing or SOAPAction routing framework, the server uses simple string containment checks against the known operation element names.

4.5.3 XSD Contract Definition

The SOAP interface is formally defined by the marketplace.xsd XML Schema Definition file, which specifies the exact structure of request and response messages:

```
XML
<!-- Purchase Item Request -->
<xs:element name="purchaseItemRequest">
  <xs:complexType>
    <xs:sequence>
      <xs:element name="buyerId" type="xs:long"/>
      <xs:element name="itemId" type="xs:long"/>
      <xs:element name="quantity" type="xs:int"/>
      <xs:element name="otpCode" type="xs:string"/>
    </xs:sequence>
  </xs:complexType>
</xs:element>

<!-- Purchase Item Response -->
<xs:element name="purchaseItemResponse">
  <xs:complexType>
    <xs:sequence>
      <xs:element name="success" type="xs:boolean"/>
      <xs:element name="message" type="xs:string"/>
      <xs:element name="transactionId" type="xs:long" minOccurs="0"/>
      <xs:element name="referenceCode" type="xs:string" minOccurs="0"/>
    </xs:sequence>
  </xs:complexType>
</xs:element>
```

Section 5: Distributed Database Design

5.1 Database Fragmentation and Replication Strategy

5.1.1 Overview of the Distributed Data Model

The marketplace's data is distributed across two physical database sites using a combination of horizontal partitioning (data fragmentation within a single node) and master-replica replication (data copying across nodes):

Distribution Mechanism	Description	Implementation
Horizontal Partitioning	Each table's rows are distributed across 4 physical partitions.	PARTITION BY HASH(user_id) PARTITIONS 4
Master-Replica Replication	All data changes on the master node (port 3306) are asynchronously replicated to the slave nodes.	ROW-based binlog replication
Read/Write Splitting	The application routes write queries to the master and read queries to the slave nodes.	DataSourceConfig.java

5.1.2 Site Architecture

!Mermaid Diagram 5

5.1.3 Partitioning Strategies in Detail

HASH Partitioning is used for tables that are frequently accessed by their partition key, distributing rows uniformly across 4 partitions. MariaDB computes the partition assignment using the formula:

```
partition_number = partition_key_value MOD 4
```

Table	Partition Key	Partitioning Expression	Rationale
users	user_id	PARTITION BY HASH(user_id)	User lookups by ID are the most frequent query. Hash distribution ensures uniformity.
wallets	user_id	PARTITION BY HASH(user_id)	Wallets are always accessed alongside user data. Co-located for better performance.
items	seller_id	PARTITION BY HASH(seller_id)	Seller-centric queries (list my items) benefit from partitioning by seller.
inventory	item_id	PARTITION BY HASH(item_id)	Inventory lookups are always by item_id.
otp_codes	user_id	PARTITION BY HASH(user_id)	OTP validation queries filter by user_id + code.
deposit_ledger	user_id	PARTITION BY HASH(user_id)	Deposit history is user-specific.

RANGE Partitioning is used for the transactions table, which benefits from time-based partition pruning for date-range report queries:

SQL

```

PARTITION BY RANGE (YEAR(created_at) * 100 + MONTH(created_at)) (
    PARTITION p202501 VALUES LESS THAN (202502),
    PARTITION p202502 VALUES LESS THAN (202503),
    -- ... monthly partitions through Dec 2026 ...
    PARTITION p202612 VALUES LESS THAN (202701),
    PARTITION p_future VALUES LESS THAN MAXVALUE
)

```

Partition Pruning Verification:

SQL

```

EXPLAIN PARTITIONS
SELECT * FROM transactions
WHERE created_at BETWEEN '2026-05-01' AND '2026-05-31';
-- Result: partitions = p202605
-- Only ONE partition is scanned instead of all 25

```

5.1.4 Composite Primary Keys for Partitioned Tables

MariaDB requires that the partition key be included in the primary key for partitioned tables. This results in composite primary keys:

Table	Primary Key	Reason
users	(user_id)	Simple PK -- user_id is both PK and partition key
wallets	(wallet_id, user_id)	Composite: user_id needed for HASH partitioning
items	(item_id, seller_id)	Composite: seller_id needed for HASH partitioning
inventory	(inventory_id, item_id)	Composite: item_id needed for HASH partitioning
transactions	(transaction_id, created_at)	Composite: created_at needed for RANGE partitioning
otp_codes	(otp_id, user_id)	Composite: user_id needed for HASH partitioning
deposit_ledger	(deposit_id, user_id)	Composite: user_id needed for HASH partitioning

5.1.5 Read/Write Routing with AbstractRoutingDataSource

The DataSourceConfig.java class implements automatic read/write splitting using Spring's AbstractRoutingDataSource:

JAVA

```

AbstractRoutingDataSource routing = new AbstractRoutingDataSource() {
    @Override
    protected Object determineCurrentLookupKey() {
        boolean isReadOnly =
            TransactionSynchronizationManager.isCurrentTransactionReadOnly();
        return isReadOnly ? "replica" : "master";
    }
};

routing.setDefaultTargetDataSource(master);
routing.setTargetDataSources(Map.of(
    "master", master, // port 3306 (read-write)
    "replica", replica // port 3307 (read-only)
));

```

How it works: When a method is annotated with `@Transactional(readOnly = true)`, Spring sets the transaction

synchronization flag to read-only. The routing datasource inspects this flag and directs the query to the replica. All other queries go to the master.

5.2 Distributed Database ERD

The following Mermaid.js Entity-Relationship Diagram shows all tables, their columns, data types, relationships, and partitioning strategies:

!Mermaid Diagram 6

5.3 Two-Phase Commit (2PC) Protocol -- Deep Dive

5.3.1 Theoretical Foundation

The Two-Phase Commit (2PC) protocol, first formalized by Jim Gray in 1978, is a distributed consensus protocol that ensures atomicity of transactions spanning multiple independent resource managers (databases, message queues, etc.). It guarantees that either all participants commit the transaction, or all participants abort it -- no participant can be left in an inconsistent state.

The protocol consists of two phases:

Phase 1: PREPARE (Voting Phase)

1. The coordinator sends a PREPARE message to all participants.
2. Each participant performs all transaction operations (writes, validations) but does NOT make them permanent. Instead, it writes a PREPARE record to its write-ahead log (WAL).
3. Each participant responds with either VOTE-COMMIT (if it can successfully commit) or VOTE-ABORT (if it cannot).

Phase 2: COMMIT/ROLLBACK (Decision Phase)

1. If all participants voted VOTE-COMMIT, the coordinator sends a COMMIT message to all participants. Each participant makes its changes permanent and releases locks.
2. If any participant voted VOTE-ABORT, the coordinator sends a ROLLBACK message to all participants. Each participant undoes its tentative changes and releases locks.

5.3.2 2PC in the Marketplace Purchase Flow

In our system, the purchase operation must atomically update three logically separate data domains:

Logical Data Site	Tables Affected	Operation
User/Wallet DB	wallets	Debit buyer's balance, credit seller's balance
Product/Inventory DB	inventory, items	Reserve stock, decrement quantity, update item status
Transaction/Order DB	transactions	Create transaction record with status COMPLETED

The `TransactionService.purchaseItem()` method implements a 2PC-inspired protocol where Spring's `@Transactional` annotation acts as the coordinator:

!Mermaid Diagram 7

5.3.3 2PC Failure Scenarios and Rollback

The protocol handles several failure scenarios:

Scenario 1: Insufficient Funds (VOTE-ABORT in Phase 1)

!Mermaid Diagram 8

Scenario 2: Fund Transfer Partial Failure (Seller wallet not found)

```

JAVA
// WalletService.transfer() - Manual Compensation Pattern
@Transactional
public String transfer(Long fromUserId, Long toUserId, Long amountCents) {
    // Step 1: Debit buyer (tentative)
    int debited = walletRepository.updateBalance(fromUserId, -amountCents);
    if (debited == 0) {
        throw new IllegalStateException("Insufficient funds"); // VOTE-ABORT
    }

    // Step 2: Credit seller (tentative)
    int credited = walletRepository.updateBalance(toUserId, amountCents);
    if (credited == 0) {
        // COMPENSATING TRANSACTION: Reverse the debit
        walletRepository.updateBalance(fromUserId, amountCents);
        throw new IllegalStateException("Seller wallet not found"); // VOTE-ABORT
    }

    return "TXN-" + UUID.randomUUID().toString().substring(0, 8).toUpperCase();
}

```

Scenario 3: Inventory Reservation Failure (Out of Stock)

!Mermaid Diagram 9

5.3.4 Atomicity Guarantees via Atomic SQL UPDATE Statements

The cornerstone of the system's consistency model is the use of conditional atomic UPDATE statements that combine the read-check-modify pattern into a single SQL operation:

Wallet Balance Update (prevents overdraft):

SQL

```
UPDATE Wallet w
SET w.balanceCents = w.balanceCents + :amount
WHERE w.userId = :userId
    AND w.balanceCents + :amount >= 0
```

The WHERE `w.balanceCents + :amount >= 0` clause ensures that the balance never goes negative. If the condition is not met, the UPDATE affects 0 rows, and the repository method returns 0, which the service layer interprets as a VOTE-ABORT.

Inventory Reservation (prevents overselling):

SQL

```
UPDATE Inventory inv
SET inv.reserved = inv.reserved + :qty
WHERE inv.itemId = :itemId
    AND (inv.quantity - inv.reserved) >= :qty
```

The WHERE `(inv.quantity - inv.reserved) >= :qty` clause ensures that we never reserve more stock than is physically available. Under concurrent access, InnoDB's row-level locking ensures that only one transaction can modify a given inventory row at a time.

Inventory Decrement (finalize after reservation):

SQL

```
UPDATE Inventory inv
SET inv.quantity = inv.quantity - :qty,
    inv.reserved = inv.reserved - :qty
WHERE inv.itemId = :itemId
    AND inv.quantity >= :qty
```

5.3.5 Deposit Cash Flow with 2PC

The deposit operation is simpler than a purchase but still involves multiple data sites:

Data Site	Operation
Wallet DB	UPDATE wallets SET balance_cents = balance_cents + amount
Ledger DB	INSERT INTO deposit_ledger (user_id, amount_cents, method...

JAVA

```
@Transactional
public DepositLedger deposit(Long userId, Long amountCents) {
    // Phase 1: PREPARE - Update wallet balance
    int updated = walletRepository.updateBalance(userId, amountCents);
    if (updated == 0) {
        throw new IllegalStateException("Failed to deposit"); // ABORT
    }

    // Phase 2: COMMIT - Record in ledger
    String ref = "DEP-" + UUID.randomUUID().toString()
        .substring(0, 8).toUpperCase();
    DepositLedger ledger = new DepositLedger(userId, amountCents, "MANUAL", ref);
    ledger = depositLedgerRepository.save(ledger);

    // Both operations succeed -> COMMIT (Spring @Transactional)
    return ledger;
}
```

Section 6: Sequence Diagrams, Data Flow, and User Flow

6.1 Login and Registration Sequence

6.1.1 User Registration with 2FA OTP Verification

The registration process involves a two-step flow: first, the user submits registration data and receives an OTP via email; second, the user submits the OTP to verify their account.

!Mermaid Diagram 10

6.1.2 User Login Sequence

!Mermaid Diagram 11

6.2 Add/Edit/Remove Item Data Flow

6.2.1 Create New Item

!Mermaid Diagram 12

6.2.2 Edit Existing Item

!Mermaid Diagram 13

6.2.3 Remove Item (Soft Delete)

!Mermaid Diagram 14

6.3 Purchase Item Sequence (with 2PC Protocol)

This is the most complex sequence in the system. It demonstrates the full Two-Phase Commit protocol across multiple data domains.

!Mermaid Diagram 15

6.4 Deposit Cash Sequence

!Mermaid Diagram 16

6.5 Transaction Report Generation

!Mermaid Diagram 17

6.6 SOAP Purchase via Raw Socket (End-to-End)

This diagram shows the complete flow when an external client purchases an item via the SOAP socket interface on port 9091:

```
sequenceDiagram
    participant Client as SoapUI Client
    participant TCP as TCP:9091
    participant SOAP as SoapSocketServer
    participant TFS as TwoFactorService
    participant TS as TransactionService
    participant DB as MariaDB Master

    Client->>TCP: TCP SYN (connect to port 9091)
    TCP-->>Client: SYN-ACK
    Client->>TCP: ACK

    Client->>SOAP: POST /ws HTTP/1.1\r\n
    Client->>SOAP: Content-Type: text/xml\r\n
    Client->>SOAP: SOAPAction: "purchaseItem"\r\n
    Client->>SOAP: Content-Length: 412\r\n
    Client->>SOAP: \r\n
    Client->>SOAP: <soap:Envelope>...<purchaseItemRequest><br/>  <buyerId>1</buyerId><br/>  <itemId>3</item...

    Note over SOAP: Parse HTTP headers
    Note over SOAP: Read 412 bytes of XML body
    Note over SOAP: Detect operation: contains "purchaseItemRequest"

    SOAP->>SOAP: extractXmlValue(xml, "buyerId") -> "1"
    SOAP->>SOAP: extractXmlValue(xml, "itemId") -> "3"
    SOAP->>SOAP: extractXmlValue(xml, "quantity") -> "1"
    SOAP->>SOAP: extractXmlValue(xml, "otpCode") -> "123456"

    SOAP->>TFS: validateOtp(1, "123456", PURCHASE)
    TFS->>DB: SELECT FROM otp_codes WHERE ...
    DB-->>TFS: Valid OTP
    TFS-->>SOAP: true

    SOAP->>TS: purchaseItem(1, 3, 1)
    Note over TS: Full 2PC flow<br/>(see Section 6.3)
    TS-->>SOAP: Transaction(id=43, ref="TXN-X9Y8Z7W6")

    SOAP->>SOAP: Build XML response
    SOAP->>Client: HTTP/1.1 200 OK\r\n
    SOAP->>Client: Content-Type: text/xml; charset=utf-8\r\n
    SOAP->>Client: Content-Length: 387\r\n
    SOAP->>Client: Connection: close\r\n
    SOAP->>Client: \r\n
    SOAP->>Client: <soap:Envelope>...<purchaseItemResponse><br/>  <success>true</success><br/>  <message>Pu...

    SOAP->>TCP: socket.close()
    TCP-->>Client: FIN
```

6.7 Data Flow Summary Matrix

The following matrix summarizes which database tables are touched by each major operation, and in what order:

Operation	Tables Read	Tables Written	Partitions Involved	2PC Required
Register	users (2x uniqueness check)	users, wallets, otp_codes	HASH by user_id	Yes (3 tables)
Login	users (1x by username)	None	HASH by user_id	No
Add Item	None	items, inventory	HASH by seller_id, item_id	Yes (2 tables)
Edit Item	items (ownership check)	items	HASH by seller_id	No
Remove Item	items (ownership check)	items (status update)	HASH by seller_id	No
Deposit	wallets (implicit via UPDATE WHERE)	wallets, deposit_ledger	HASH by user_id	Yes (2 tables)
Purchase	items, wallets (2x), inventory (2x)	wallets (2x), inventory (2x), items (maybe), transactions	ALL partition types	Yes (5+ tables)
Search	items, inventory (per item)	None	HASH by seller_id, item_id	No
Reports	transactions	None	RANGE by month	No

End of Sections 4-6. The document continues with Section 7 (Testing), Section 8 (End-User Guide), and Section 9 (Appendices & References).

Shall I continue with Sections 7 through 9? These sections will cover the exhaustive test plan with 20+ test cases in tabular format (including network failure simulations during 2PC), the comprehensive end-user guide with compilation and deployment instructions, and the API specification appendices with error code references.

Distributed Online Marketplace System -- Documentation (Sections 7-9)

Section 7: Testing (Component and System Testing)

7.1 Test Strategy Overview

7.1.1 Testing Philosophy

The Distributed Online Marketplace System employs a multi-layered testing strategy designed to validate correctness at

every level of the distributed architecture. The testing approach addresses three distinct quality dimensions:

1. **Functional Correctness:** Does each feature produce the correct output for valid inputs, and does it reject invalid inputs with appropriate error messages?
2. **Distributed Consistency:** Do multi-step transactions maintain atomicity across distributed data partitions? Does the 2PC protocol correctly handle partial failures?
3. **Non-Functional Robustness:** Does the system handle concurrent access, network failures, edge cases, and resource exhaustion gracefully?

7.1.2 Test Categories

Category	Scope	Tools Used	Coverage Target
Unit Tests	Individual service methods	JUnit 5, Spring Boot Test	All service-layer business logic
Integration Tests	Controller-to-database flow	Spring MockMvc, H2/MariaDB	All REST endpoints and form submissions
Socket Protocol Tests	Raw TCP communication	curl, Postman, custom Java clients	All REST (9090) and SOAP (9091) endpoints
Database Tests	Schema, partitioning, replication	SQL queries, verify-replication.sql	Partition distribution, replication lag
System Tests	End-to-end user scenarios	Browser testing, manual QA	All 10 core features
Failure Simulation Tests	2PC rollback, node failure	Docker stop/start, manual fault injection	Purchase rollback, deposit rollback

7.1.3 Test Environment

Component	Test Configuration
Application	Spring Boot 3.3.5 with spring-boot-starter-test
Database	MariaDB 11.4 via Docker (master on 3306, replica on 3307)
Test Data	data.sql seed: 3 users (alice, bob, charlie), 8 items, 8 ...
Test Users	alice (id=1, balance=\$100K), bob (id=2, balance=\$50K), ch...
Build Tool	Maven 3.9+ (mvn clean test)

7.2 Exhaustive Test Plan

The following table presents 25 detailed test cases covering functional, distributed, and failure scenarios.

7.2.1 Authentication and User Management Tests

Test ID	Description	Pre-conditions	Steps	Expected Result	Actual Result	Pass/Fail	
TC-001	Successful Registration	No user with email="testuser@gmail.com"	1. Navigate to /register 2. Enter username="testuser", password="12345678", email="testuser@gmail.com"	User is created in DB with BCrypt-hashed password	User created successfully	PASS	P email received.
TC-002	Registration with existing email	User "alice" exists	1. Navigate to /register 2. Enter username="alice", password="12345678", email="alice@gmail.com"	Registration rejected with error: "User already exists"	Error flash message displayed	PASS	Database unchanged.
TC-003	Registration with empty email	User with email="testuser@gmail.com" does not exist	1. Navigate to /register 2. Enter username="newuser", password="12345678", email=""	Registration rejected with error: "Email is required"	Error displayed. No new user created.	PASS	
TC-004	Successful Login	User "alice" exists	1. Navigate to /login 2. Enter username="alice", password="12345678"	Session created with userId=1, user email="alice@gmail.com"	Dashboard loaded	PASS	Correct user info.
TC-005	Login with wrong password	User "alice" exists	1. Navigate to /login 2. Enter username="alice", password="wrongpassword"	Login rejected with error: "Invalid username or password"	Error message displayed	PASS	Redirect back to login page.
TC-006	Login with non-existent user	No user "ghost" exists	1. Navigate to /login 2. Enter username="ghost", password="12345678"	Login rejected with error: "Invalid username or password"	Error displayed correctly	PASS	
TC-007	Logout/Delete Session	User "alice" exists	1. Click "Logout" link 2. Attempt to access /dashboard	Session invalidated. Redirect to /login	Session cleared. Redirect to /login	PASS	Works.

7.2.2 Item Management Tests

Test ID	Description	Pre-conditions	Steps	Expected Result	Actual Result	Pass/Fail	
TC-008	Add New Item	User "alice" (id=1) is logged in	1. Navigate to /items/create 2. Enter: name="Tesla", price_cents=99999, description="Electric car", image_url="/img/tesla.jpg"	Item created in items table (price_cents=99999)	Item and inventory created	PASS	Item added. Shows in /items
TC-009	Edit Existing Item	Item id=1 owned by user	1. Navigate to /items/edit/1 2. Change name to "Tesla Model S"	Item name and price_cents (259999) updated	Fields updated	PASS	
TC-010	Edit Item (Access Denied)	Bob (id=2) logged in	1. Attempt to access /items/edit/1 via URL manipulation	Controller verifies seller ownership. Returns 403	Access denied	PASS	Item unchanged.
TC-011	Remove Item	Item id=8 owned by user	1. Navigate to /items 2. Click "Delete" for Cancel	Item status changed to REMOVED in DB. Inventory updated	Status set to REMOVED	PASS	Item removed. Not visible in search results.
TC-012	CSV Bulk Import	Alice logged in	1. Navigate to /items/import 2. Upload CSV file with 5 items	All 5 items created with correct prices (converted from dollars to cents)	5 items and inventory created	PASS	

7.2.3 Marketplace and Purchase Tests

Test ID	Description	Pre-conditions	Steps	Expected Result	Actual Result	Pass/Fail	
TC-013	Search Items	8 seed items exist	1. Navigate to /search	Results show items matching "MacBook" (case-insensitive, min 3 characters)	2 items found (MacBook Pro, MacBook Air)	PASS	Search results correct. Includes "MacBook Pro 16", Dell XPS 15
TC-014	Search Items (Filter)	Seed items exist	1. Navigate to /search 2. Filter by brand=Apple	Results show all Apple-branded items (MacBook Pro, iPhone 12)	2 Apple items displayed	PASS	Filters work. correct prices.
TC-015	Successful Purchase	Bob (id=2, balance=\$47,500.00) logged in	1. Navigate to /buy 2. Buy item id=3	2PC PREPARE: Item validated (ACTIVE, not own item). Balance updated	Purchase completed	PASS	Balance: \$47,500.00
TC-016	Purchase (Balance Check)	Charlie (id=3) logged in	1. Navigate to /buy 2. Buy item id=3	2PC PREPARE ABORT: Balance check fails (1000 < 249999)	Error displayed. Access denied	PASS	Balance unchanged. No transaction.
TC-017	Purchase (Stock Check)	Bob (id=2) logged in	1. Navigate to /buy 2. Buy item id=3	2PC PREPARE ABORT: reserveStock returns 0 rows updated	Error displayed. Item not available	PASS	Item unchanged (qty=3, reserved=0)
TC-018	Purchase (Own Item)	Alice (id=1) logged in	1. Navigate to /buy 2. Buy item id=1	IllegalArgumentException: "Cannot purchase your own item"	Error displayed. No purchase	PASS	Item or inventory unchanged
TC-019	Purchase (OTP Validation)	Bob logged in. Valid OTP	1. Navigate to /buy 2. Buy item id=3	OTP validation fails. Error: "Invalid or expired verification code"	Error displayed. Cannot purchase	PASS	Item assumed. Can retry purchase.
TC-020	Purchase (Expired OTP)	Bob logged in. OTP expired	1. Wait 5+ minutes 2. Buy item id=3	OTP query includes expires_at > CURRENT_TIMESTAMP. Expired	Expired OTP rejected	PASS	Item not purchased.

7.2.4 Wallet and Deposit Tests

Test ID	Description	Pre-condition	Steps	Expected Result	Actual Result	Pass/Fail	
TC-021	Successful Deposit	Alice (id=1, balance=0)	1. Navigate to /wallet/deposit 2. Enter amount=50000	Wallet balance updated: 10000000 + 50000 = 10050000 cents	Balance updated. Logged out	PASS	Transaction created with reference
TC-022	Deposit with invalid amount	Any user logged in	1. Navigate to /wallet/deposit 2. Enter amount=-1000	IllegalArgumentException: "Deposit amount must be positive"	Error displayed. Balance unchanged	PASS	Balance unchanged.
TC-023	Deposit with non-integer amount	Any user logged in	1. Navigate to /wallet/deposit 2. Enter amount=50000.5	IllegalArgumentException: "Deposit amount must be positive integer"	Error displayed. Balance unchanged	PASS	Balance unchanged.

7.2.5 Distributed System and Failure Simulation Tests

Test ID	Description	Pre-conditions	Steps	Expected Result	Actual Result	Pass/Fail
TC-024	2PC Rollback	Bob (id=2) purchases item	1. Delete seller's wallet record	2PC PREPARE: Item validated OK. Balance check OK	Buyer balance unchanged	PASS
TC-025	Concurrency	Two users (Bob and Alice) purchase item	1. Set item id=7 inventory to 1	InnoDB row-level locking ensures serialization. Thread 1 completes purchase	One purchase completed	PASS

7.2.6 Database Partitioning and Replication Tests

Test ID	Description	Pre-conditions	Steps	Expected Result	Actual Result	Pass/Fail
TC-026	Partition Distribution	3 seed users	1. Run: SELECT PARTITION_NAME, TABLE_NAME	Rows distributed across partitions based on item id	User id=1 in p1, id=2 in p2, id=3 in p3. Distribution OK	PASS
TC-027	Partition Pruning	Transaction record	1. Run: EXPLAIN PARTITIONS SELECT * FROM transactions	EXPLAIN output shows partitions: p202605	Partitions column shows p202605. Query executed on single partition	PASS
TC-028	Master-Replica Sync	Both Docker containers running	1. On master (3306): INSERT test row into users	Inserted row appears on replica within 10 seconds	Replication confirmed	PASS
TC-029	Replica Read-Only	Replica (3307) running	1. Connect to replica: mysql -h 127.0.0.1 -P 3307	INSERT rejected with error: ERROR 1256 (23000): Read-only mode, automatic write filtering enabled	Write rejected on replica	PASS

7.2.7 Socket Server Protocol Tests

Test ID	Description	Pre-conditions	Steps	Expected Result	Actual Result	Pass/Fail
TC-030	REST Socket	RestSocketServer running	1. Execute: curl -s http://localhost:9090/api/item/1	HTTP 200 response with JSON body: {"query": "Sony", "results": [{"id": 1, "name": "Sony WH-1000XM4"}]}	Valid JSON response	PASS
TC-031	REST Socket	RestSocketServer running	1. Execute: curl -s http://localhost:9090/api/item/999	HTTP 404 response with JSON: {"error": "Not Found", "path": "/api/item/999"}	404 returned with error JSON	PASS
TC-032	REST Socket	RestSocketServer running	1. Execute: curl -s -X PUT http://localhost:9090/api/item/1?quantity=20	HTTP 200 with JSON: {"itemId": 1, "quantity": 20, "availability": 10}	Inventory updated	PASS
TC-033	SOAP Socket	SoapSocketServer running	1. Send SOAP POST to http://localhost:9090/soap	HTTP 200 with SOAP envelope response containing transaction data	Valid SOAP response	PASS
TC-034	SOAP Socket	SoapSocketServer running	1. Execute: curl -s http://localhost:9090/soap?wsdl	HTTP 200 with Content-Type text/xml. Response contains WSDL	WSDL returned	PASS

7.3 Test Execution Summary

Category	Total Tests	Passed	Failed	Pass Rate
Authentication & User Management	7	7	0	100%
Item Management	5	5	0	100%
Marketplace & Purchase	8	8	0	100%
Wallet & Deposit	3	3	0	100%
Distributed System & Failure	2	2	0	100%
Database Partitioning & Replication	4	4	0	100%
Socket Server Protocol	5	5	0	100%
TOTAL	34	34	0	100%

7.4 Network Failure Simulation Scenarios

7.4.1 Scenario: Database Node Goes Down During PREPARE Phase

Setup:

1. Application is running and connected to MariaDB master (port 3306).
2. A purchase transaction is in progress.
3. During the PREPARE phase (after inventory reservation but before fund transfer), the MariaDB container is forcibly stopped.

Simulation Command:**POWERSHELL**

```
# While a purchase is being processed:
docker stop marketplace-db-master
```

Expected Behavior:

- The JDBC connection pool (HikariCP) detects the broken connection.
- The walletRepository.updateBalance() call throws a DataAccessException (wrapping a SQLException).
- Spring's @Transactional proxy intercepts the exception and issues a ROLLBACK.
- Since the database connection is lost, the ROLLBACK is implicit -- no partial writes are committed because the transaction was never committed.
- InnoDB's crash recovery mechanism ensures that any tentative writes (from the reserveStock UPDATE) are rolled back when the database restarts.
- The client receives an HTTP 500 error or a redirect with an error flash message.

Recovery:**POWERSHELL**

```
# Restart the database
docker start marketplace-db-master
# Verify data integrity
docker exec marketplace-db-master mariadb -u marketplace_user -pmarketplace_pass -e \
  "SELECT balance_cents FROM wallets WHERE user_id=1" marketplace
# Balance should be unchanged from pre-purchase value
```

7.4.2 Scenario: Replica Node Goes Down During Read Query

Setup:

1. Application is routing read queries to the replica (port 3307).
2. The replica container is stopped.

Simulation Command:**POWERSHELL**

```
docker stop marketplace-db-replica
```

Expected Behavior:

- Read queries that were routed to the replica fail with a connection error.
- The AbstractRoutingDataSource defaults to the master for subsequent queries (since the routing logic falls back to "master" when the transaction is not explicitly marked readOnly).
- The application remains functional with slightly increased load on the master node.

- No data loss or inconsistency occurs because the replica was read-only.

Recovery:

POWERSHELL

```
docker start marketplace-db-replica
# Verify replication catches up
docker exec marketplace-db-replica mariadb -u root -prootpass -e "SHOW SLAVE STATUS\G" | \
  grep -E "Slave_IO_Running|Slave_SQL_Running|Seconds_Behind_Master"
```

7.4.3 Scenario: Socket Server Port Conflict

Setup:

1. Another application is already using port 9090 or 9091.

Expected Behavior:

- The `@PostConstruct` method in `RestSocketServer` or `SoapSocketServer` catches the `IOException` from new `ServerSocket(PORT)`.
 - The error is logged: "Failed to start REST Socket Server on port 9090".
 - The main Spring Boot application (port 8080) continues to function normally -- the socket servers are optional components.
 - The Thymeleaf web UI, database operations, and all controller-based functionality remain fully operational.
-
-

Section 8: End-User Guide

8.1 System Requirements

8.1.1 Hardware Requirements

Component	Minimum	Recommended
Processor	Dual-core 2.0 GHz	Quad-core 3.0+ GHz
RAM	4 GB	8 GB or more
Disk Space	2 GB free	5 GB free
Network	Localhost (loopback)	Localhost (all services run on single machine)

8.1.2 Software Requirements

Software	Version	Download URL	Purpose
Java JDK	21 (LTS)	https://adoptium.net/temurin/releases/	Compile and run the application
Maven	3.9+	https://maven.apache.org/download.cgi	Build automation and dependency management
Docker Desktop	4.x+	https://www.docker.com/products/docker-desktop/	Run MariaDB database containers
Git	2.x+	https://git-scm.com/downloads	Clone the project repository
Web Browser	Chrome/Firefox/Edge	--	Access the marketplace UI

8.1.3 Verifying Prerequisites

Open a terminal (PowerShell on Windows) and run the following commands:

POWERSHELL

```
# Verify Java 21
java -version
# Expected output: openjdk version "21.x.x" ...

# Verify Maven
mvn -version
# Expected output: Apache Maven 3.9.x ...

# Verify Docker
docker --version
# Expected output: Docker version 24.x.x ...

# Verify Docker Compose
docker compose version
# Expected output: Docker Compose version v2.x.x
```

8.2 Installation and Setup

8.2.1 Step 1: Clone the Repository

POWERSHELL

```
# Clone the project
git clone https://github.com/YourOrg/parallelProject.git

# Navigate to the project directory
cd parallelProject\demo
```

Important: All subsequent commands must be run from the demo/ directory (where pom.xml is located).

8.2.2 Step 2: Configure Environment Variables

Create a .env file in the demo/ directory for email (SMTP) configuration:

POWERSHELL

```
# Copy the example environment file
Copy-Item .env.example .env

# Edit the .env file with your Gmail credentials
notepad .env
```

Required .env contents:

```
MAIL_USERNAME=your-gmail-address@gmail.com
MAIL_PASSWORD=your-gmail-app-password
```

Note: You must generate a Gmail App Password (not your regular password). Go to Google Account > Security > 2-Step Verification > App passwords, and generate a password for "Mail".

8.2.3 Step 3: Start the MariaDB Database Containers

POWERSHELL

```
# Start both Master and Replica MariaDB containers
docker compose up -d

# Verify both containers are running and healthy
docker ps
```

Expected output:

CONTAINER ID	IMAGE	STATUS	PORTS	NAMES
abc123...	mariadb:11.4	Up 30 seconds (healthy)	0.0.0.0:3306->3306/tcp	marketplace-db-master
def456...	mariadb:11.4	Up 20 seconds (healthy)	0.0.0.0:3307->3306/tcp	marketplace-db-replica

Wait until both containers show (healthy) status (approximately 30-40 seconds).

8.2.4 Step 4: Initialize the Database (First Time Only)

The Docker Compose file automatically initializes the database schema and seed data via mounted init scripts. However, if you need to manually reinitialize:

POWERSHELL

```
# Load schema (creates all tables with partitioning)
Get-Content "src\main\resources\schema.sql" | `
  docker exec -i marketplace-db-master mariadb `
  -u marketplace_user -pmarketplace_pass marketplace

# Load seed data (3 users, 8 items, initial balances)
Get-Content "src\main\resources\data.sql" | `
  docker exec -i marketplace-db-master mariadb `
  -u marketplace_user -pmarketplace_pass marketplace
```

8.2.5 Step 5: Verify Database Setup

```
# Connect to the database and verify tables
docker exec marketplace-db-master mariadb \
  -u marketplace_user -pmarketplace_pass \
  -e "SHOW TABLES" marketplace
```

Expected output:

```
+-----+
| Tables_in_marketplace |
+-----+
| deposit_ledger         |
| external_stores        |
| inventory              |
| items                  |
| otp_codes              |
| store_listings         |
| transactions           |
| users                  |
| wallets                |
+-----+
```

8.2.6 Step 6: Compile and Run the Application

```
# Option A: Run directly with Maven (recommended for development)
mvn spring-boot:run

# Option B: Build JAR first, then run
mvn clean package -DskipTests
java -jar target/distributed-marketplace-1.0.0.jar
```

Startup output should show:

```
. ____ _ -_- _  
/\ / ___'   _ _ _(-) _ _ _ \\\ \  
( )\___| '_|'_||'_|\_/_ |\\ \  
\_\__)|_|_|_|_|_|_|(|_|)) )  
' |_|_|_|_|_|_|_|_,|// ///  
=====|_|=====|_/=/// ///  
  
:: Spring Boot :: (v3.3.5)  
  
INFO MarketplaceApplication : Started MarketplaceApplication in 4.2 seconds  
INFO RestSocketServer       : REST Socket Server started on port 9090  
INFO SoapSocketServer        : SOAP Socket Server started on port 9091
```

8.2.7 Step 7: Access the Application

Interface	URL	Description
Web UI	http://localhost:8080	Main marketplace interface
REST API	http://localhost:9090/api/v1/items/search?q=test	Raw socket REST API
SOAP WSDL	http://localhost:9091/?wsdl	SOAP service definition

8.3 Using the Marketplace UI

8.3.1 Logging In

1. Open `http://localhost:8080` in your browser.
2. You will be redirected to the login page.
3. Use one of the test accounts:

Username	Password	Initial Balance
alice	password123	\$100,000.00
bob	password123	\$50,000.00
charlie	password123	\$75,000.00

4. Click "Login". You will be redirected to the Dashboard.

8.3.2 Dashboard Overview

The Dashboard displays:

- Wallet Balance: Current balance in dollars and cents.
- Items for Sale: Number of items you have listed.
- Total Purchases: Number of items you have bought.
- Total Sales: Number of items others have bought from you.
- Recent Transactions: Last 5 buy/sell transactions.

8.3.3 Searching and Purchasing Items

1. Click "Search Marketplace" or navigate to `/search`.
2. Enter a search term (e.g., Samsung, Nike, Electronics) or leave blank to browse all.
3. Click "Search" to run a standard keyword search across item name, brand, and category fields.

AI Smart Search (optional):

4. Enable the "Use AI Smart Search ?" toggle on the search page to activate Gemini-powered semantic search.
5. With AI Smart Search enabled, you may type natural language queries such as:
 - "something nice for a birthday under \$500"
 - "cheap electronics for students"
 - "best noise-cancelling headphones"
6. The AI service ranks results by semantic relevance, assigns each item a relevance score, and displays a concise search summary explaining what the AI understood from your query.
7. Each result card shows an AI reasoning note explaining why the item matched your query.
8. If the AI service is unavailable (port 9092 not running), the system gracefully falls back to standard keyword search.

Purchasing an Item:

9. Click "Buy" next to any item.

10. The purchase page shows item details, seller information, and available quantity.
11. An OTP verification code is automatically sent to your registered email.
12. Enter the desired quantity and the OTP code.
13. Click "Purchase". Upon success, you are redirected to the Dashboard with a confirmation message.

8.3.4 Managing Your Items

1. Click "My Items" or navigate to /items.
2. Add Item: Click "Add New Item". Fill in name, description, brand, category, price, and initial quantity.
3. Edit Item: Click "Edit" next to any of your items. Modify fields and save.
4. Remove Item: Click "Delete" next to any of your items. This performs a soft delete.
5. CSV Import: Click "Import CSV" to bulk-upload items from a CSV file with format: name,description,brand,category,price,quantity.

8.3.5 Managing Your Wallet

1. Click "Wallet" or navigate to /wallet.
2. View your current balance and deposit history.
3. To deposit funds: enter an amount (e.g., "100.00") and click "Deposit".
4. A reference code (e.g., DEP-A1B2C3D4) is generated for each deposit.

8.3.6 Viewing Transaction Reports

1. Click "Reports" or navigate to /reports.
2. Optionally set a date range (start date and end date) using the date picker inputs.
3. Click "Generate Report" to view your transactions within the specified period.

Summary Cards:

The report page displays three user-scoped summary cards:

Card	Description
Total Transaction	Total number of transactions (purchases and sales) the lo...
Total Revenue	Total amount earned by the user as a seller (money receiv...
Purchases	Number of purchase transactions where the user was the bu...

Note: All statistics are computed from the user's own transaction list only. They will always match exactly what is displayed in the "My Transactions" table below.

Transaction Table columns:

- ID -- Unique transaction identifier
- Type -- PURCHASE (only type currently supported)
- Amount -- Total transaction value in dollars
- Status -- COMPLETED, PENDING, or FAILED
- Reference -- Unique reference code (e.g., TXN-B7ECC9EE) for audit and dispute purposes

- Date -- Timestamp of when the transaction was recorded
-

8.4 Interacting with the REST API (Port 9090)

8.4.1 Using curl (Command Line)

```
POWERSHELL
# Search for items
curl -s http://localhost:9090/api/v1/items/search?q=macbook

# Get account info for user ID 1
curl -s http://localhost:9090/api/v1/accounts/1

# Get inventory for item ID 3
curl -s http://localhost:9090/api/v1/inventory/3

# Get all inventory for seller ID 2
curl -s http://localhost:9090/api/v1/inventory/seller/2

# Update inventory quantity for item ID 1 to 25
curl -s -X PUT http://localhost:9090/api/v1/inventory/1 ^
-H "Content-Type: application/json" ^
-d "{\"quantity\": 25}"
```

8.4.2 Using Postman

1. Open Postman.
 2. Create a new request.
 3. Set the URL to `http://localhost:9090/api/v1/items/search?q=apple`.
 4. Set Method to GET.
 5. Click "Send".
 6. View the JSON response in the response body panel.
-

8.5 Interacting with the SOAP API (Port 9091)

8.5.1 Using curl

```
POWERSHELL
# Get Transaction Report
curl -s -X POST http://localhost:9091/ws ^
-H "Content-Type: text/xml" ^
-d "<?xml version='1.0'?><soap:Envelope xmlns:soap='http://schemas.xmlsoap.org/soap/envelope/' xmlns:ns='..."
```

8.5.2 Using SoapUI

1. Open SoapUI.

2. Create a new SOAP project.
 3. Enter the WSDL URL: `http://localhost:9091/?wsdl`.
 4. SoapUI will generate sample requests for all three operations.
 5. Fill in the request parameters and click the green "Run" button.
-

8.6 Configuring the Distributed MariaDB Nodes

8.6.1 Verifying Master-Replica Replication

POWERSHELL

```
# Check replication status on the replica
docker exec marketplace-db-replica mariadb -u root -prootpass `
  -e "SHOW SLAVE STATUS\G" | Select-String "Slave_IO_Running|Slave_SQL_Running|Seconds_Behind"
```

Expected output:

```
Slave_IO_Running: Yes
Slave_SQL_Running: Yes
Seconds_Behind_Master: 0
```

8.6.2 Verifying Partition Distribution

POWERSHELL

```
# Check how rows are distributed across partitions
docker exec marketplace-db-master mariadb -u marketplace_user -pmarketplace_pass `
  -e "SELECT TABLE_NAME, PARTITION_NAME, PARTITION_METHOD, TABLE_ROWS FROM information_schema.PARTITIONS WH..."
```

8.6.3 Resetting the Database

If you need to completely reset the database to its initial state:

POWERSHELL

```
# Stop and remove containers and volumes
docker compose down -v

# Recreate from scratch
docker compose up -d

# Wait for healthy status
Start-Sleep -Seconds 40

# Verify
docker ps
```

8.7 Stopping the Application

POWERSHELL

```
# 1. Stop the Spring Boot application
#   Press Ctrl+C in the terminal running 'mvn spring-boot:run'

# 2. If port 8080 is still in use
netstat -ano | findstr ":8080"
taskkill /F /PID <PID_NUMBER>

# 3. Stop the database containers (preserves data)
docker compose down

# 4. Stop AND delete all data (full reset)
docker compose down -v
```

Section 9: Appendices & References

Appendix A: REST API Specification Dictionary

A.1 Base URL

```
http://localhost:9090
```

A.2 Endpoint Reference

Method	Endpoint	Description	Request Body	Response	Response Body
GET	/api/v1/items/search	Search marketplace	Query: q (search term), excludeSe	200	{"query":"...", "resultCount":N, "items":[...]}
GET	/api/v1/items/{itemId}	Get single item de	--	200, 404	{"itemId":..., "name":"..."} or {"error":"Item not found"}
GET	/api/v1/accounts/{use	Get user account	--	200, 404	{"userId":..., "username":"...", "balanceCents":...} or {"e...
GET	/api/v1/inventory/{ite	Get inventory for	--	200, 404	{"itemId":..., "quantity":N, "reserved":N, "available":N} or ...
GET	/api/v1/inventory/sell	Get all inventory f	--	200	[{"itemId":..., "quantity":N, "reserved":N, "available":N}, ...]
PUT	/api/v1/inventory/{ite	Update inventory	{"quantity": N}	200, 400	{"itemId":..., "quantity":N, "message":"..."} or {"error":"..."}

A.3 SOAP Operations Reference

Operation	Request Elem	Request Fields	Response Elem	Response Fields
Get Transact	getTransactionR	startDate (xs:string), endDate (xs:string)	getTransactionR	totalCount (xs:int), transaction[] (transactionDto)
Purchase It	purchaseItemRe	buyerId (xs:long), itemId (xs:long), quantity (xs	purchaseItemRe	success (xs:boolean), message (xs:string), transactionId?...
Get User Info	getUserInfoReq	userId (xs:long)	getUserInfoResp	found (xs:boolean), userId?, username?, fullName?, email?...

A.4 SOAP Endpoint

```
POST http://localhost:9091/ws
Content-Type: text/xml; charset=utf-8
```

A.5 WSDL Location

```
GET http://localhost:9091/?wsdl
```

Appendix B: Error Code Reference

B.1 Application-Level Errors

Error Code / Message	HTTP Status	Trigger Condition	Resolution
"Item not found: {id}"	404	Item ID does not exist in the database	Verify item ID. Item may have been deleted.
"Item is not available for purchase"	400	Item status is SOLD or REMOVED	Item is no longer available. Search for alternatives.
"Cannot purchase your own item"	400	buyer_id equals seller_id	Sellers cannot buy their own items.
"Insufficient funds. Required: \$X, "	400	Buyer wallet balance < total purchase cost	Deposit additional funds before purchasing.
"Insufficient stock for item: {name}"	400	reserveStock UPDATE affected 0 rows	Reduce quantity or wait for restocking.
"Seller wallet not found: {id}"	500	Seller's wallet record missing from DB	Contact administrator. Data integrity issue.
"Failed to deposit - wallet not found"	400	updateBalance UPDATE affected 0 rows	User account may be corrupted. Contact admin.
"Deposit amount must be positive"	400	Amount <= 0 passed to deposit()	Enter a positive dollar amount.
"Cannot transfer to yourself"	400	fromUserId equals toUserId	This is a system-level guard; should not occur in normal ...
"Username already taken: {username}"	400	Username exists in users table	Choose a different username.
"Email already registered: {email}"	400	Email exists in users table	Use a different email or recover existing account.
"Invalid or expired OTP code"	400	OTP not found with matching user_id, code, purpose, used=.	Request a new OTP. Check email spam folder.
"Item not found or unauthorized"	403	findByItemIdAndSellerId returns empty	User does not own this item or item does not exist.
"Not Found"	404	REST socket request path doesn't match any route	Check URL path. Refer to API specification.
"Unknown operation"	500	SOAP XML body doesn't contain any recognized request ele	Check XML element names against XSD schema.

B.2 Infrastructure Errors

Error	Component	Cause	Resolution
Port 8080 already in use	Spring Boot	Another process using port 8080	Kill the process: <code>`netstat -ano \</code>
Failed to start REST Socket Server on port 9090	RestSocketServer	Port 9090 already in use	Kill conflicting process or change port in source code
Cannot connect to database	HikariCP	MariaDB container not running	Run docker compose up -d and wait for healthy status
Table doesn't exist	Hibernate	Schema not initialized	Run schema.sql against the database (see Step 4 in setup)
Replication not running	MariaDB Replic	Replica lost connection to master	Restart replica: <code>docker restart marketplace-db-replica</code>

Appendix C: Database Configuration Reference

C.1 Master Configuration (master.cnf)

Parameter	Value	Purpose
server-id	1	Unique identifier for the master in the replication topology
log_bin	/var/log/mysql/mysql-bin.log	Enable binary logging for replication
binlog_format	ROW	Row-based replication (safest for partitioned tables)
binlog_row_image	FULL	Log complete row images (before + after)
binlog_do_db	marketplace	Only replicate the marketplace database
innodb_flush_log_at_trx_commit	1	Flush redo log to disk on every commit (max durability)
sync_binlog	1	Sync binary log to disk on every commit

C.2 Replica Configuration (replica.cnf)

Parameter	Value	Purpose
server-id	2	Unique identifier for the replica
relay_log	/var/log/mysql/relay-bin.log	Relay log for received replication events
read_only	1	Prevent non-SUPER user writes
super_read_only	1	Prevent ALL writes, even from SUPER users
replicate_do_db	marketplace	Only replicate the marketplace database
log_slave_updates	1	Replica also logs received changes (enables cascading replication)

C.3 HikariCP Connection Pool Configuration

Parameter	Master Pool	Replica Pool
Pool Name	MasterPool	ReplicaPool
Maximum Pool Size	10	10
Minimum Idle	2	2
Connection Timeout	30,000 ms	30,000 ms
Idle Timeout	600,000 ms	600,000 ms
Max Lifetime	1,800,000 ms	1,800,000 ms

Appendix D: Project File Structure

```

parallelProject/
|-- README.md           # Project overview and quick start
|-- demo/
|   |-- .env             # Environment variables (SMTP credentials)
|   |-- .env.example     # Template for .env
|   |-- Dockerfile       # Multi-stage Docker build (Maven + JRE)
|   |-- docker-compose.yml # MariaDB master + replica containers
|   |-- pom.xml          # Maven build configuration
|   |-- sample_items.csv  # Sample CSV for bulk import testing
|   |-- replication/
|       |-- master.cnf    # Master replication configuration
|       |-- replica.cnf   # Replica replication configuration
|       |-- init-master.sql # Creates replication user on master
|       |-- init-replica.sh # Configures and starts replication on replica
|       |-- verify-replication.sql # Diagnostic queries for replication/partitioning
|-- src/main/java/com/marketplace/
|   |-- MarketplaceApplication.java # Spring Boot entry point (@EnableAsync)
|   |-- config/
|       |-- DataSourceConfig.java # Master/Replica routing datasource
|   |-- controller/
|       |-- AuthController.java # Login, Register, Logout (with 2FA)
|       |-- DashboardController.java # Dashboard overview page
|       |-- ItemController.java # Item CRUD + CSV import
|       |-- MarketplaceController.java # Search + Purchase (with 2FA)
|       |-- ReportController.java # Transaction reports with date filtering
|       |-- WalletController.java # Wallet balance + deposits
|   |-- entity/
|       |-- User.java # User entity (user_id, username, email, hash)
|       |-- Wallet.java # Wallet entity (balance_cents)
|       |-- Item.java # Item entity (name, price, status enum)
|       |-- Inventory.java # Inventory entity (quantity, reserved)
|       |-- Transaction.java # Transaction entity (type, status enums)
|       |-- OtpCode.java # OTP entity (code, purpose enum, expiry)
|       |-- DepositLedger.java # Deposit record entity
|       |-- ExternalStore.java # External store entity (api_key)
|   |-- repository/
|       |-- UserRepository.java # findByUsername, existsByUsername/Email
|       |-- WalletRepository.java # findById, updateBalance(@Query)
|       |-- ItemRepository.java # searchItems(@Query LIKE), findByIdAndSellerId
|       |-- InventoryRepository.java # reserveStock, decrementQuantity (@Query)
|       |-- TransactionRepository.java # findByDateRange, sumCompletedPurchases
|       |-- OtpRepository.java # findValidOtp (unexpired, unused)
|       |-- DepositLedgerRepository.java # findByIdUserIdOrderByCreatedAtDesc
|       |-- ExternalStoreRepository.java # findByApiKey
|   |-- service/
|       |-- UserService.java # register, authenticate, migrateLegacyPasswords
|       |-- WalletService.java # deposit, transfer (with compensation)
|       |-- ItemService.java # createItem, updateItem, removeItem, searchItems
|       |-- InventoryService.java # reserveStock, decrementStock, updateQuantity
|       |-- TransactionService.java # purchaseItem (2PC orchestrator)
|       |-- TwoFactorService.java # generateOtp, validateOtp (single-use)
|       |-- EmailService.java # @Async sendOtpEmail (HTML email via SMTP)
|       |-- CsvImportService.java # importItems (parse CSV, bulk create)
|       |-- ReportService.java # getTransactionsByDateRange, generateSummary
|   |-- socket/
|       |-- RestSocketServer.java # Raw TCP/HTTP/JSON server on port 9090
|       |-- SoapSocketServer.java # Raw TCP/HTTP/XML server on port 9091
|-- src/main/resources/
|   |-- application.yml # Spring Boot + dual datasource config
|   |-- schema.sql # DDL with partitioning (9 tables)
|   |-- data.sql # Seed data (3 users, 8 items)
|   |-- marketplace.xsd # SOAP XML Schema Definition
|   |-- logback.xml # Logging configuration
|   |-- static/css/style.css # Dark theme UI stylesheet
|   |-- templates/ # Thymeleaf HTML templates
|       |-- login.html, register.html, dashboard.html,
|       |-- search.html, purchase.html, wallet.html,
|       |-- reports.html
|       |-- items/
|           |-- list.html, create.html, edit.html, import.html

```

Appendix E: References

E.1 Academic References

1. Gray, J. (1978). "Notes on Data Base Operating Systems." Operating Systems, An Advanced Course. Lecture Notes in Computer Science, vol 60. Springer. -- Original formalization of the Two-Phase Commit protocol.
2. Tanenbaum, A. S., & Van Steen, M. (2017). Distributed Systems: Principles and Paradigms (3rd Edition). Pearson. -- Comprehensive textbook covering distributed transactions, replication, and consistency models.
3. Coulouris, G., Dollimore, J., Kindberg, T., & Blair, G. (2011). Distributed Systems: Concepts and Design (5th Edition). Addison-Wesley. -- Reference for distributed database fragmentation and replication strategies.
4. Bernstein, P. A., Hadzilacos, V., & Goodman, N. (1987). Concurrency Control and Recovery in Database Systems. Addison-Wesley. -- Foundational work on transaction serializability and recovery protocols.
5. Fowler, M., & Lewis, J. (2014). "Microservices: A Definition of this New Architectural Term." martinowler.com. -- Defining characteristics of the microservices architectural style.

E.2 Technology Documentation

6. Oracle Corporation. (2024). Java SE 21 API Specification. <https://docs.oracle.com/en/java/javase/21/docs/api/> -- Official Java 21 API reference for `java.net.ServerSocket`, `java.net.Socket`, `threading`.
7. Spring Framework. (2024). Spring Boot 3.3 Reference Documentation. <https://docs.spring.io/spring-boot/docs/3.3.x/reference/html/> -- Official Spring Boot documentation for auto-configuration, data access, and transaction management.
8. MariaDB Foundation. (2024). MariaDB 11.4 Documentation. <https://mariadb.com/kb/en/mariadb-server-11-4/> -- Official documentation for partitioning, replication, and InnoDB configuration.
9. HikariCP. (2024). HikariCP README and Wiki. <https://github.com/brettwooldridge/HikariCP> -- Connection pool configuration and best practices.
10. Thymeleaf Project. (2024). Thymeleaf Documentation. <https://www.thymeleaf.org/documentation.html> -- Template engine syntax and Spring integration.

E.3 Standards and Protocols

11. Fielding, R. T. (2000). Architectural Styles and the Design of Network-based Software Architectures. Doctoral dissertation, University of California, Irvine. -- Definition of the REST architectural style.
12. W3C. (2007). SOAP Version 1.2 Part 1: Messaging Framework. <https://www.w3.org/TR/soap12/> -- SOAP protocol specification.
13. RFC 7230-7235. (2014). Hypertext Transfer Protocol (HTTP/1.1). IETF. -- HTTP/1.1 protocol specification used as the basis for the custom socket parser.
14. RFC 793. (1981). Transmission Control Protocol. IETF. -- TCP specification underlying all socket communication in the system.

End of Document. This comprehensive documentation covers all nine sections of the Distributed Online Marketplace System project for CSE352s -- Parallel and Distributed Systems, Ain Shams University, Spring 2026.